

---

# Ego-Conditioned Pedestrian Body Motion Generation via Latent Diffusion with Temporal Cross-Attention

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 We introduce EgoPed, a system that generates realistic 3D full-body pedestrian  
2 motions conditioned on vehicle ego trajectory observations from autonomous driv-  
3 ing data. Existing pedestrian motion synthesis methods rely on text descriptions or  
4 action class labels, ignoring the driving context that fundamentally shapes pedes-  
5 trian behavior — how a person walks, stops, or reacts is strongly influenced by  
6 nearby vehicle speed and proximity. We condition a Motion Latent Diffusion  
7 (MLD) model on vehicle ego odometry via cross-attention, using a frozen ego  
8 encoder that maps the full 196-timestep ego trajectory sequence to per-timestep  
9 conditioning tokens attending to the denoiser at every diffusion step. Experiments  
10 on a combined dataset built from nuScenes, Waymo, and recordings from our  
11 sensor-equipped research vehicle (AVA) demonstrate a **49% FID improvement**  
12 ( $6.603 \rightarrow 3.392$ ) over the strongest pooled-conditioning baseline, and an abla-  
13 tion study reveals that the frozen encoder is critical: unfreezing causes a 34%  
14 MultiModality collapse ( $3.956 \rightarrow 2.618$ ) as the encoder over-specializes for dis-  
15 criminative per-condition features. Our best system achieves  $\text{FID} = 3.392 \pm 0.179$   
16 with  $\text{R-Precision@1} = 0.548 \pm 0.002$  at  $\text{CFG} = 10$ , providing a strong founda-  
17 tion for generating behavior-consistent pedestrian motions for AV simulation and  
18 safety-critical scenario generation.

## 19 1 Introduction

20 Generating realistic pedestrian body motions is a critical capability for autonomous vehicle (AV)  
21 simulation. Simulation pipelines require diverse, plausible pedestrian behaviors — walking, turning,  
22 stopping, reacting to traffic — to stress-test perception and planning systems before deployment.  
23 While substantial progress has been made in text-conditioned 3D human motion generation Chen et al.  
24 [2023], Tevet et al. [2023], Zhang et al. [2024], Guo et al. [2024], these systems generate motions  
25 that are context-agnostic: a pedestrian crossing a street looks identical regardless of whether a vehicle  
26 is approaching at 50 km/h or stationary.

27 In reality, pedestrian behavior is strongly coupled to the surrounding vehicle context. The ego vehicle’s  
28 speed, trajectory, and proximity all influence whether a pedestrian will cross, wait, accelerate, or turn.  
29 This *ego-conditioned* setting — generating full 3D body motion conditioned on vehicle ego trajectory  
30 — is largely unexplored. Existing work in this space either generates 2D trajectory distributions  
31 only Mao et al. [2023], or uses ego context as scene conditioning without 3D body synthesis Feng  
32 et al. [2024].

33 This paper addresses ego-conditioned *full-body* pedestrian motion synthesis. We build on Motion  
34 Latent Diffusion (MLD) Chen et al. [2023], which compresses 263D HumanML3D body motion  
35 into a compact latent space and denoises within it, but replace the text conditioning with vehicle

ego trajectory. The key architectural question is: *how should the ego trajectory be injected into the denoiser?*

We study two families of conditioning mechanisms:

- **Pooled conditioning (H2):** The ego trajectory is encoded by a transformer encoder and mean-pooled to a single 256D token, which is concatenated alongside the time embedding in the denoiser’s self-attention. This gives the denoiser a coarse summary of the ego trajectory but no temporal resolution.
- **Temporal cross-attention conditioning (H4):** The ego trajectory is encoded to a full 196-token sequence, and the denoiser uses cross-attention ( $z$  queries ego as K/V at each denoising step). This exposes every denoising step to the full ego trajectory, enabling temporally fine-grained conditioning.

Our main finding is that temporal cross-attention dramatically outperforms pooled conditioning. H4 achieves **FID = 3.392**, a 49% improvement over the best H2 result (FID = 6.603), with the same denoiser capacity — and its best checkpoint occurs *earlier* in training (epoch 3399 vs. 4399), so the improvement does not come from additional compute. The cross-attention denoiser can attend to specific moments in the ego trajectory rather than a coarse average, producing motion that more closely follows the ground truth distribution.

A secondary but equally important finding concerns the ego encoder’s training regime. H4 uses a *frozen* ego encoder — pretrained to align mean-pooled ego trajectories with VAE motion latents, then fixed during diffusion training. We ablate this choice with H6, which unfreezes the encoder and allows it to co-adapt with the denoiser. H6 consistently *underperforms* H4 across all evaluated checkpoints: FID worsens by 50%, R-Precision@1 drops by 36%, and MultiModality collapses by 34%. We show that this occurs because the unfrozen encoder specializes to discriminative, per-condition representations that over-constrain generation diversity — consistent with the frozen-encoder paradigm in image generation (frozen CLIP in Stable Diffusion Rombach et al. [2022], adapters on frozen backbones in IP-Adapter Ye et al. [2023] and ControlNet Zhang et al. [2023b]).

Our contributions are:

1. The first system to generate full 3D HumanML3D body motions conditioned on vehicle ego odometry, trained on AVA, nuScenes, and Waymo data.
2. A demonstration that temporal cross-attention over the full ego sequence achieves 49% FID improvement over single-token pooled conditioning, with non-overlapping confidence intervals across 3 evaluation replications.
3. An ablation showing that frozen ego encoders are critical for generative diversity in ego-conditioned motion: unfreezing causes MultiModality collapse and worsens all metrics.
4. A design recommendation: frozen ego encoder + cross-attention denoiser (H4) is the correct architecture for this task.

## 2 Related Work

**Ego-conditioned and context-conditioned motion generation.** The closest works to ours generate 2D pedestrian trajectories conditioned on ego-vehicle or scene context. LED Mao et al. [2023] uses latent-space diffusion for stochastic trajectory prediction conditioned on past observations, demonstrating that latent-space diffusion substantially outperforms raw-space approaches for trajectory generation. UniTraj Feng et al. [2024] trains a unified trajectory-prediction model across multiple AV datasets (nuScenes, Argoverse 2, Waymo Open Motion), demonstrating the value of multi-dataset training for agent behavior modeling — though for 2D trajectory prediction rather than 3D body synthesis. Our work differs in generating 263D full-body motion in the HumanML3D representation, enabling realistic body pose, gait, and foot contact — not just root position.

**Interaction-aware and retrieval-augmented motion generation.** ReMoDiffuse Zhang et al. [2023c] augments motion diffusion with retrieved nearest-neighbour motions cross-attended into the denoiser, an architectural pattern related to ours where external sequences condition the denoiser via cross-attention. InterGen Liang et al. [2024] jointly generates two-person interactions with mutual




Figure 1: EgoPed architecture. The ego trajectory is encoded by a *frozen* transformer encoder into 196 per-timestep conditioning tokens. At every denoising step, the noisy latent  $z_t$  queries the full ego token sequence via cross-attention (trans\_dec denoiser). Diffusion runs in the latent space of a frozen HumanML3D VAE; only the denoiser is trained.

cross-attention. Our setting is a degenerate version where one agent (the ego vehicle) has a fixed, observed trajectory — enabling a simpler asymmetric architecture where ego trajectory serves as the conditioning K/V.

**Latent diffusion for human motion.** MLD Chen et al. [2023] is our backbone: it diffuses in the latent space of a pre-trained HumanML3D VAE, using a cross-attention transformer denoiser conditioned on CLIP text embeddings. We retain the VAE, denoiser architecture, evaluation metrics, and DDIM sampling from MLD, replacing text conditioning with ego trajectory. MDM Tevet et al. [2023] operates in raw motion space and conditions on text, action labels, or trajectory waypoints. T2M-GPT Zhang et al. [2023a] uses discrete VQ-VAE tokenization — less suitable for continuous ego trajectory signals. MoMask Guo et al. [2024] achieves state-of-the-art FID on HumanML3D with masked transformer generation; we note that our metrics are computed in the same evaluation protocol, though on a different ego-conditioned dataset split.

**Classifier-free guidance.** CFG Ho and Salimans [2022] trains conditional and unconditional models jointly, enabling quality-diversity tradeoff at inference. We find that CFG interacts strongly with conditioning granularity: pooled conditioning (H2) shows monotonically increasing FID with CFG scale, while temporal cross-attention (H4) is nearly CFG-insensitive. This suggests that richer conditioning naturally prevents the over-constraint that degrades FID at high CFG.

**Frozen encoders in generative models.** A recurring pattern in conditional image generation is that frozen encoders preserve generative diversity better than fine-tuned ones. Stable Diffusion Rombach et al. [2022] freezes CLIP during diffusion training; IP-Adapter Ye et al. [2023] and ControlNet Zhang et al. [2023b] add conditioning adapters on frozen diffusion backbones. Conversely, approaches that fine-tune the generative stack end-to-end, such as DreamBooth Ruiz et al. [2023], are known to trade diversity for condition fidelity and require explicit prior-preservation regularization to counteract it. Our H6 ablation provides direct empirical evidence for the same phenomenon in motion generation: unfreezing the ego encoder degrades generation quality and collapses per-condition diversity.

### 3 Method

#### 3.1 Problem Formulation

Given a vehicle ego trajectory  $\mathbf{e} \in \mathbb{R}^{T \times 2}$  (2D position over  $T = 196$  timesteps), we aim to generate a plausible 3D pedestrian body motion  $\mathbf{x} \in \mathbb{R}^{T \times 263}$  in the HumanML3D representation Guo et al. [2022]. The 263D representation encodes root velocity, root angular velocity, joint positions, velocities, and foot contact labels, fully parameterizing a SMPL-like body in motion.

We model the conditional distribution  $p(\mathbf{x} | \mathbf{e})$  using a latent diffusion model that denoises within a pre-trained VAE’s latent space.

#### 3.2 Motion VAE

We use the MLD VAE Chen et al. [2023], a Transformer encoder-decoder that compresses  $T \times 263$  motion to a latent  $\mathbf{z} \in \mathbb{R}^{L \times 256}$  with  $L = 4$  latent tokens ( $L$  is a hyperparameter; we find  $L = 4$  significantly outperforms  $L = 8$ , see Section 5). The VAE is pre-trained on the HumanML3D dataset and held fixed throughout diffusion training.

#### 3.3 Ego Encoder

The ego trajectory  $\mathbf{e} \in \mathbb{R}^{196 \times 2}$  is encoded by a lightweight transformer encoder:

$$\mathbf{c} = \text{EgoEncoder}(\mathbf{e}) \in \mathbb{R}^{196 \times 256}. \quad (1)$$

The encoder consists of a linear projection from 2D to 256D, followed by sinusoidal positional encodings, and  $N = 4$  transformer encoder layers with 4 heads. Unlike the pooled variant (H2), which applies mean-pooling to produce a single token  $\mathbf{c} \in \mathbb{R}^{1 \times 256}$ , our encoder preserves the full temporal sequence.

**Ego encoder pretraining.** Before diffusion training, the ego encoder is pretrained with a regression-based alignment objective that distills the paired motion’s VAE latent into the ego encoding. Specifically, we align the mean-pooled ego encoding  $\bar{\mathbf{c}} = \text{MeanPool}(\mathbf{c}) \in \mathbb{R}^{256}$  with the mean-pooled VAE posterior mean  $\bar{\mathbf{z}} \in \mathbb{R}^{256}$  (computed from the paired motion) using a combined MSE and cosine objective:

$$\mathcal{L}_{\text{pretrain}} = \|\bar{\mathbf{c}} - \bar{\mathbf{z}}\|^2 + \lambda_{\cos}(1 - \cos(\bar{\mathbf{c}}, \bar{\mathbf{z}})), \quad (2)$$

where the cosine term encourages directional alignment in addition to magnitude. After pretraining, the encoder achieves validation MSE  $\approx 1.92$ , matching the pooled variant (which uses an additional 2-layer projection head without improvement), confirming that the projection head adds no alignment capacity.

**Frozen encoder (key design choice).** During diffusion training, the ego encoder parameters are *frozen*. This preserves the encoder’s learned alignment while ensuring the generated motion distribution is not over-constrained to particular ego conditions. We ablate this choice in Section 5.

#### 3.4 Diffusion Denoiser

We use the MLD denoiser with `trans_dec` architecture: a transformer *decoder* where the noisy latent  $\mathbf{z}_t$  forms the queries and the ego conditioning tokens  $\mathbf{c}$  form the keys and values. The denoiser predicts the noise component,

$$\hat{\epsilon} = \epsilon_{\theta}(\mathbf{z}_t, t, \mathbf{c}) = \text{TransformerDecoder}(\mathbf{z}_t; [\mathbf{t}_{\text{emb}}, \mathbf{c}]), \quad (3)$$

where the time step  $t$  is embedded and prepended to the K/V sequence alongside  $\mathbf{c}$ , and the sampler computes  $\mathbf{z}_{t-1}$  from  $\hat{\epsilon}$ . Cross-attention makes full-sequence conditioning affordable: attention cost is  $\mathcal{O}(L \times T)$ , whereas concatenating all  $T = 196$  ego tokens into the denoiser’s *self*-attention would cost  $\mathcal{O}((L + T)^2)$ . Relative to the pooled baseline (self-attention over only  $L + 2$  tokens), the added cost is modest — and in exchange, every denoising step gains temporal resolution over the full ego trajectory.

At each denoising step, the 4 latent queries attend to all 196 ego tokens, enabling temporally fine-grained conditioning: the denoiser can attend to specific moments in the ego trajectory (e.g., a braking event or sharp turn) rather than a global average.

### 155 3.5 Classifier-Free Guidance

156 We train with classifier-free guidance Ho and Salimans [2022]: for 10% of training samples, the  
 157 input ego trajectory is zeroed before encoding, so the encoder’s response to the all-zero trajectory  
 158 serves as the unconditional embedding (the same construction is used for the unconditional branch at  
 159 inference). At inference, the guided score is:

$$\hat{\epsilon}_t = \epsilon_\theta(\mathbf{z}_t, \emptyset) + w \cdot (\epsilon_\theta(\mathbf{z}_t, \mathbf{c}) - \epsilon_\theta(\mathbf{z}_t, \emptyset)), \quad (4)$$

160 with guidance scale  $w$ . We find  $w = 10$  optimal for H4 (Section 4).

### 161 3.6 Training Objective

162 The denoiser is trained with the standard DDPM noise-prediction objective:

$$\mathcal{L} = \mathbb{E}_{\mathbf{z}_0, \epsilon, t} [\|\epsilon - \epsilon_\theta(\mathbf{z}_t, t, \mathbf{c})\|^2]. \quad (5)$$

163 We use DDIM Song et al. [2021] for fast inference with 50 denoising steps.

## 164 4 Experiments

### 165 4.1 Dataset

166 As there are no datasets available for the problem investigated in this work, we generated a dataset  
 167 for ego-conditioned motion generation from both existing autonomous driving datasets and from  
 168 scenes recorded with our sensor-equipped research vehicle AVA. The vehicle is equipped with seven  
 169 128-layer LiDARs and eight RGB cameras for full 360° sensor coverage around the vehicle. **Figure ??**  
 170 **shows the vehicle and a detailed sensor setup.** To obtain scenes with meaningful interaction between  
 171 the vehicle and a pedestrian, we staged **NUMBER** highly interactive scenarios. **An example scene of**  
 172 **a pedestrian waving a vehicle through can be seen in Figure ??.** The recorded data is brought into  
 173 nuScenes format and objects in the scene are annotated with 3D bounding boxes and orientation  
 174 labels using MS3D Tsai et al. [2024, 2023].

175 With the data in nuScenes format fully annotated, we extract 3D human poses in SMPL Loper et al.  
 176 [2015] format using an adapted version of OmniRe Chen et al. [2025] from all three data sources:  
 177 AVA, nuScenes, and Waymo Open Dataset. We estimate poses for every pedestrian instance from all  
 178 available camera images, using the projected 3D bounding box to localize the person in each view.  
 179 Per-instance pose estimates from different cameras are then stitched into a single continuous sequence  
 180 using the extrinsic calibration of the camera rigs. Finally, the SMPL pose sequences are converted  
 181 to the 263D HumanML3D motion representation Guo et al. [2022]. In addition to the pedestrian  
 182 pose sequences, we extract the synchronized ego-vehicle trajectory from each scene’s odometry. A  
 183 single automated pipeline processes every (scene, pedestrian) pair end-to-end: it loads the annotated  
 184 nuScenes-format scene, extracts and stitches the per-instance SMPL body sequence, resamples both  
 185 the body motion and the ego odometry to a common 20 fps grid, and expresses both streams in  
 186 a *pedestrian-centric* coordinate frame (translation- and yaw-normalized to the pedestrian) so that  
 187 conditioning is invariant to the absolute world pose of the scene. Each resulting sample is stored as  
 188 a self-contained record containing the pedestrian’s 22 body-joint positions ( $\text{ped\_in\_ped\_frame} \in \mathbb{R}^{T \times 22 \times 3}$ ),  
 189 the ego trajectory in the same frame ( $\text{ego\_in\_ped\_frame} \in \mathbb{R}^{T \times 3}$ ), and the derived  
 190 263D HumanML3D motion vector ( $\text{vectors\_263} \in \mathbb{R}^{T \times 263}$ ), where the native sequence length  $T$   
 191 varies per scene and is windowed to 196 frames at training time (Section 4.2 details the cropping).

192 Using all three data sources we obtain **11,954** paired ego–pedestrian motion samples in total, split  
 193 by scene into 9,540 training and 2,414 validation samples with no scene overlap between the two  
 194 splits. Per-source counts are given in Table 1. Waymo and nuScenes dominate the sample count,  
 195 while the AVA recordings contribute a small but deliberately high-interaction subset of staged  
 196 vehicle–pedestrian encounters.

197 Each sample is a paired sequence of ego odometry  $\mathbf{e} \in \mathbb{R}^{196 \times 2}$  and pedestrian body motion  $\mathbf{x} \in \mathbb{R}^{196 \times 263}$ ,  
 198 both expressed in the pedestrian-centric frame.

199 **Interaction scoring.** Rather than hard-filtering samples, we assign each scenario a scalar *interaction*  
 200 *score* that quantifies how strongly the pedestrian and ego vehicle interact, and use it both to bias

Table 1: Dataset composition. Each sample is a paired sequence of ego odometry ( $196 \times 2$ ) and pedestrian body motion ( $196 \times 263$ ), extracted from AVA, nuScenes, and Waymo. Train/validation splits are disjoint at the scene level; the validation split is the evaluation set used for all reported metrics.

| Source             | Train | Validation | Total  |
|--------------------|-------|------------|--------|
| AVA (ours)         | 27    | 7          | 34     |
| nuScenes           | 3,637 | 943        | 4,580  |
| Waymo Open Dataset | 5,876 | 1,464      | 7,340  |
| Total              | 9,540 | 2,414      | 11,954 |

training-time sampling and to guide cropping. Let  $\mathbf{r}_t, \mathbf{g}_t \in \mathbb{R}^2$  denote the ground-plane (XZ) positions of the pedestrian pelvis and the ego vehicle at frame  $t$ . The score combines three factors:

$$s = \underbrace{\sum_t \|\mathbf{r}_{t+1} - \mathbf{r}_t\|_2}_{\text{pedestrian travel}} \cdot \underbrace{\frac{1}{T} \sum_t \mathbb{I}[\|\mathbf{g}_t - \mathbf{r}_t\|_2 < 5 \text{ m}]}_{\text{proximity fraction}} \cdot \left( 1 + \frac{1}{180^\circ} \sum_t |\Delta\phi_t| \right), \quad (6)$$

where  $\phi_t = \text{atan2}(\mathbf{r}_t - \mathbf{g}_t)$  is the pedestrian’s bearing relative to the ego vehicle and  $\Delta\phi_t$  is the (wrap-corrected) inter-frame change in bearing. Intuitively, a scenario scores highly when the pedestrian moves substantially, spends many frames within 5 m of the vehicle, and sweeps a large relative bearing (e.g., crossing in front of the vehicle).

**Interaction-weighted sampling.** During diffusion (and VAE) training we draw samples with a `WeightedRandomSampler` whose weights are the interaction scores  $s$  (floored at 0.01 so low-interaction scenarios are still occasionally seen, then normalized to sum to the dataset size). This concentrates gradient updates on the ego-relevant portion of the data distribution without discarding any samples. Evaluation uses uniform (unweighted) sampling.

**Interaction-aware cropping.** To bound the generation problem to a fixed horizon, we crop each scenario to  $T = 196$  timesteps. For scenarios longer than 196 frames, rather than taking a random window we centre the crop on the *closest-approach frame*  $t^* = \arg \min_t \|\mathbf{g}_t - \mathbf{r}_t\|_2$ , then add uniform jitter of  $\pm 25\%$  of the window length to the centre so the moment of closest approach is not memorized at a fixed position; the resulting window is clamped to lie within the sequence. Ego and motion streams are cropped with a shared index so they stay temporally aligned. Scenarios shorter than 196 frames are zero-padded at the end and their true length is recorded for masking.

## 4.2 Evaluation Protocol

We report four metrics, following the HumanML3D evaluation protocol Guo et al. [2022]:

- **FID** (Fréchet Inception Distance): computes the feature-space distance between the generated and real motion distributions, using the pretrained `t2m_moveencoder` (stride-2 Conv1D) and `t2m_motionencoder` (BiGRU) from Guo et al. [2022]. Lower is better.
- **R-Precision@1**: top-1 retrieval accuracy of the ego condition from batches of 32 candidates (1 positive + 31 negatives), using cosine similarity between the mean-pooled VAE latent (from generated motion) and the ego encoder output. Higher is better.
- **MultiModality (MM)**: average pairwise L2 distance between 10 generated samples for the same ego condition, measuring diversity per condition. Higher is better.
- **Diversity**: average pairwise L2 distance in feature space across all generated samples, measuring population-level diversity. Values *closer to the ground-truth diversity* (5.496 on our test set) are better.

All reported values are the mean over **3 independent replications** with 95% confidence intervals. Note that FID and Diversity use a fixed, pretrained evaluator shared by all systems, so they are directly comparable across architectures. R-Precision, in contrast, retrieves with each model’s

Table 2: Reference baselines bounding the metric space (val split; same evaluator, normalization, and split lists as Table 3). All rows are mean  $\pm$  95% CI over 3 replications (the unconditional row resamples the diffusion model; the non-learned rows are reseeded end-to-end). Retrieval attains near-zero FID by copying real motions but is non-generative; trajectory+kinematic shows that following the root trajectory with a static body is catastrophic. Ego conditioning improves EgoPed’s FID by 34% over its own unconditional prior ( $5.18 \rightarrow 3.39$ ) — and even the unconditional H4 prior outperforms conditionally generated H2 motions, underscoring the architectural advantage of the cross-attention denoiser.

| Method                            | FID $\downarrow$  | Diversity | R@1 $\uparrow$     |
|-----------------------------------|-------------------|-----------|--------------------|
| Retrieval (NN ego)*               | 0.062 $\pm$ 0.004 | 5.34      | –                  |
| Trajectory+kinematic (mean)       | 49.60 $\pm$ 0.087 | 0.46      | –                  |
| Trajectory+kinematic (oracle)     | 47.67 $\pm$ 0.085 | 0.72      | –                  |
| VAE latent interpolation          | 7.91 $\pm$ 0.033  | 5.82      | –                  |
| Unconditional EgoPed (ego zeroed) | 5.18 $\pm$ 0.209  | 5.96      | 0.031 <sup>†</sup> |
| H2 (Pooled)                       | 6.603             | 5.78      | 0.671              |
| H4 (EgoPed, ours)                 | <b>3.392</b>      | 5.79      | 0.548              |

\*Non-generative upper bound: copies real training motions, so low FID reflects memorization rather than conditioned synthesis. <sup>†</sup>Exactly chance level ( $1/32 = 0.031$ ), confirming the R-Precision protocol: without ego information, retrieval is random.

own ego encoder embedding; cross-architecture R@1 comparisons (e.g., pooled vs. full-sequence encoders) should therefore be read as indicative rather than exact, as the retrieval space itself differs.

### 4.3 Baselines

- **H2 (Pooled, our baseline):** EgoEncoderPooled (transformer + mean-pool  $\rightarrow$  single 256D token), trans\_enc denoiser (concatenated self-attention over  $[\mathbf{z}, \mathbf{c}_{\text{pool}}, \mathbf{t}]$ ). Best at epoch 4399, CFG = 5.
- **H3 (Latent-8):** Same as H2 but with latent dimension  $L = 8$  (larger VAE). Best at epoch 4399, CFG = 5.

We additionally report four non-learned or reference baselines (Table 2) that bound the metric space. **Unconditional MLD** runs our H4 model (epoch 3399) with every ego trajectory zeroed; because CFG training dropout uses the same zeroed-trajectory construction, both guidance branches receive identical embeddings and guidance cancels exactly, yielding the model’s pure unconditional prior and isolating the contribution of ego conditioning. **Retrieval** returns, for each test ego trajectory, the training motion whose ego trajectory is the nearest neighbour (32-point resampled L2). **VAE interpolation** decodes a random latent interpolation of two training motions. **Trajectory+kinematic** replaces all body features with the training mean, keeping either the ground-truth root velocity (*oracle*) or nothing (*mean*), simulating a trajectory-prediction-then-fitting pipeline with no learned body model.

### 4.4 Main Results

Table 3 shows the definitive evaluation results. H4 (EgoPed) achieves  $\text{FID} = 3.392 \pm 0.179$  at epoch 3399 and  $\text{CFG} = 10$ , representing a **49% improvement** over the best baseline (H2:  $\text{FID} = 6.603 \pm 0.067$  at  $\text{CFG} = 5$ ). The improvement is significant: the confidence intervals are non-overlapping (H4 CI: [3.213, 3.571]; H2 CI: [6.536, 6.670]).

Interestingly, R-Precision@1 for H4 (0.548) is *lower* than H2’s baseline (0.671). We attribute this to three compounding factors: (1) H4 generates more diverse motions per ego condition (MM: 3.956 vs. 3.503), spreading the generated distribution further in feature space and naturally reducing retrieval accuracy; (2) the R@1 metric measures alignment between generated VAE latents and ego encoder outputs — as the denoiser learns richer dynamics, this alignment metric underestimates the true conditioning fidelity; (3) R@1 retrieves with each model’s own ego encoder (Section 4.2), so the pooled encoder (trained with a projection head specifically for this alignment) and the full-sequence encoder define retrieval tasks of different intrinsic difficulty. The large FID improvement confirms

Table 3: Main evaluation results. Each method is reported at its own optimal CFG (sweep in Table 4). All values are mean  $\pm$  95% CI over 3 replications. Best values in **bold**; for Diversity, values closer to the ground-truth reference are better. H4 = EgoPed (ours).

| Method                          | Epoch | CFG | FID $\downarrow$         | R@1 $\uparrow$    | MM $\uparrow$            | Diversity |
|---------------------------------|-------|-----|--------------------------|-------------------|--------------------------|-----------|
| H2 (Pooled, baseline)           | 4399  | 5   | 6.603 $\pm$ 0.067        | 0.671 $\pm$ 0.003 | 3.503                    | 5.779     |
| H3 (Latent-8 VAE)               | 4399  | 5   | 7.563 $\pm$ 0.078        | 0.676 $\pm$ 0.003 | 3.245 $\pm$ 0.386        | 5.842     |
| H4 (Ours, FREEZE_EGO=True)      | 3399  | 10  | <b>3.392</b> $\pm$ 0.179 | 0.548 $\pm$ 0.002 | <b>3.956</b> $\pm$ 0.406 | 5.794     |
| H6 (Ablation, FREEZE_EGO=False) | 3399  | 10  | 5.079 $\pm$ 0.034        | 0.352 $\pm$ 0.009 | 2.618 $\pm$ 0.309        | 6.064     |
| Ground truth (reference)        | —     | —   | —                        | —                 | —                        | 5.496     |

Table 4: CFG sweep at epoch 3399. H4 FID is nearly flat across guidance scales; H2 degrades steeply.

| Method      | CFG | FID $\downarrow$         | R@1 $\uparrow$           | MM $\uparrow$ |
|-------------|-----|--------------------------|--------------------------|---------------|
| H2 (Pooled) | 5   | 6.603 $\pm$ 0.067        | 0.671 $\pm$ 0.003        | 3.503         |
| H2 (Pooled) | 10  | 6.963 $\pm$ 0.035        | <b>0.767</b> $\pm$ 0.006 | 3.031         |
| H2 (Pooled) | 15  | 7.400 $\pm$ 0.016        | <b>0.797</b> $\pm$ 0.005 | 2.724         |
| H4 (Ours)   | 5   | 3.842 $\pm$ 0.108        | 0.506 $\pm$ 0.009        | 4.141         |
| H4 (Ours)   | 10  | <b>3.392</b> $\pm$ 0.179 | 0.548 $\pm$ 0.002        | 3.956         |
| H4 (Ours)   | 15  | 3.547 $\pm$ 0.122        | 0.536 $\pm$ 0.003        | 3.910         |

that H4 produces a generated distribution *much* closer to the ground truth, even as retrieval accuracy decreases.

## 4.5 CFG Scale Analysis

We sweep  $\text{CFG} \in \{5, 10, 15\}$  for H4 at epoch 3399 (Table 4). A striking finding: H4 FID is nearly flat across CFG scales (3.842  $\rightarrow$  3.617), while H2 shows steep monotonic FID degradation at higher CFG (6.603  $\rightarrow$  7.400). This indicates that temporal cross-attention naturally provides a rich, overdetermined conditioning signal — the denoiser attends to 196 ego tokens, so higher guidance amplifies an already high-fidelity signal rather than forcing mode collapse onto a coarse pooled token.

## 5 Ablation Study

### 5.1 Latent Dimension: 4 vs. 8 (H3)

To test whether a larger VAE latent space improves generation quality, we train a latent-8 VAE and ego-conditioned diffusion model (H3). As shown in Table 3, H3 achieves FID = 7.563 — **14.5% worse** than H2’s 6.603 at the same epoch and CFG. R-Precision is essentially unchanged (H3: 0.676 vs. H2: 0.671), confirming that the bottleneck for conditioning quality is the ego encoder, not the latent space.

The mechanism: the fixed-capacity denoiser must now model an 8D latent distribution instead of 4D, a strictly harder task, without any additional denoiser parameters. The marginal diversity improvement (Diversity: 5.842 vs. 5.779, +1.1%) does not compensate. **Lesson:** increasing latent dimension is a poor lever for FID improvement without scaling the denoiser capacity proportionally.

### 5.2 Encoder Freeze vs. Unfreeze (H6)

The most important ablation is whether to freeze the ego encoder during diffusion training. H6 uses an identical architecture to H4 but with FREEZE\_EGO=False, allowing the encoder to co-adapt with the denoiser via backpropagation through the diffusion loss.

Figure 3 shows the H6 training trajectory (FID and R@1 vs. epoch). Two behaviors are notable:






Figure 2: Qualitative results from H4 (epoch 3399, CFG=10) on three ego-conditioned samples from our test set. Each row shows five timesteps of the generated 3D body motion (HumanML3D, rendered as a rigged skeleton on a small ground plane) alongside a bird’s-eye-view trajectory plot (rightmost column) showing the predicted pedestrian root trajectory (blue), ground-truth root trajectory (orange), and ego-vehicle trajectory (red). The body motion exhibits realistic gait, foot contact, and orientation changes that respond to the ego context.

- **FID/R@1 tradeoff:** As training progresses, R@1 improves (encoder specializes discriminatively) but FID degrades (generation quality decreases). The two optima do *not* coincide: training-time FID is best at epoch 2099 (4.789), while R@1 continues rising long after FID has begun to deteriorate.
- **R@1 saturation far below H4:** After epoch  $\approx 3299$ , training-time R@1 saturates, oscillating in the 0.37–0.42 band through epoch 4999 (maximum 0.417 at epoch 4599) while FID remains between 5.1 and 6.3. Neither metric ever approaches H4’s operating point (FID 3.392, R@1 0.548) at any point in 5000 epochs.

Definitive evaluations (3 replications, CFG = 10) confirm the picture:

Table 5: H6 definitive evaluation at three checkpoints. H4 reference included for comparison.

| System                             | FID ↓                              | R@1 ↑                              | MM ↑         | Diversity |
|------------------------------------|------------------------------------|------------------------------------|--------------|-----------|
| H6 epoch=2099 (best train FID)     | 4.952 $\pm$ 0.088                  | 0.263 $\pm$ 0.009                  | 2.455        | 5.957     |
| H6 epoch=3299 (best seg-1 R@1)     | 5.311 $\pm$ 0.060                  | 0.348 $\pm$ 0.004                  | 2.581        | 6.034     |
| H6 epoch=3399 (matched to H4 best) | 5.079 $\pm$ 0.034                  | 0.352 $\pm$ 0.009                  | 2.618        | 6.064     |
| <b>H4 epoch=3399 (FREEZE=True)</b> | <b>3.392 <math>\pm</math>0.179</b> | <b>0.548 <math>\pm</math>0.002</b> | <b>3.956</b> | 5.794     |

H6 is worse than H4 on *all metrics at all evaluated checkpoints*. Later checkpoints cannot change this conclusion: even H6’s best training-time R@1 in the saturation band (0.417 at epoch 4599, with training-time FID 6.339) remains far below H4’s definitive R@1 of 0.548 — and training-time metrics consistently *overestimate* the definitive 3-replication values (e.g., 0.398  $\rightarrow$  0.348 at epoch 3299). The MultiModality collapse is particularly striking: H6 MM = 2.618 vs. H4 MM = 3.956 (34% lower). This indicates near-deterministic generation per condition — the ego encoder has learned highly discriminative features for each trajectory, collapsing the generated motion distribution to a narrow region.

**Mechanism.** When the encoder is frozen (H4), it maps ego trajectories to general-purpose representations aligned with the VAE latent space. The denoiser learns to query these fixed representations and generate diverse, plausible motions for each condition. When the encoder is unfrozen (H6), backpropagation through the cross-attention loss pushes the encoder toward representations that *minimize the denoising loss* — essentially learning discriminative features that uniquely identify each trajectory. This over-specialization constrains generation: the denoiser outputs similar samples for each unique ego representation, collapsing MultiModality.




Figure 3: H6 training trajectory (FREEZE\_EGO=False): training-time validation FID (left axis, red) and R@1 (right axis, blue) vs. epoch, over the full 5000-epoch run. FID is best at epoch 2099 and degrades as R@1 rises; after epoch  $\approx 3299$ , R@1 saturates in the 0.37–0.42 band while FID oscillates between 5.1 and 6.3. Dashed reference lines show H4’s best operating point (FID=3.392, R@1=0.548 at epoch=3399, CFG=10). H6 never approaches H4’s performance on either metric.

314 This pattern mirrors the frozen-encoder paradigm in image generation Rombach et al. [2022], Ye et al.  
 315 [2023], Zhang et al. [2023b]: frozen conditioning encoders preserve general-purpose representations  
 316 and allow the generative model to learn diverse ways to satisfy the conditioning signal.

### 317 5.3 Checkpoint Selection

318 A non-obvious finding (relevant to practitioners) is that FID does not monotonically improve with  
 319 training. H2’s epoch trajectory shows FID = 6.603 at epoch 4399, FID = 8.400 at epoch 4599  
 320 (regression), and FID = 7.510 at epoch 4999 (partial recovery) — despite flat training loss throughout.  
 321 H4 similarly peaks at epoch 3399 (definitive FID = 3.392) and degrades steadily thereafter (Fig. 4).  
 322 **Lesson:** the final checkpoint is not the best checkpoint. Periodic FID validation during training is  
 323 essential; early stopping on FID is critical.

## 324 6 Discussion

325 **Architecture recommendation.** The experimental evidence strongly favors the H4 architecture:  
 326 frozen ego encoder + cross-attention trans\_dec denoiser. Future work seeking to improve conditioning  
 327 fidelity (R@1) should explore adapter-style approaches — fine-tuning a lightweight adapter on the  
 328 frozen encoder while keeping the backbone fixed — rather than end-to-end unfreezing.

329 **R-Precision@1 and the conditioning-diversity tradeoff.** H4’s R@1 (0.548) is lower than H2’s  
 330 (0.671) despite dramatically better FID. This seemingly paradoxical result reflects a genuine tradeoff:  
 331 richer conditioning via temporal cross-attention enables the denoiser to generate *more diverse* motions  
 332 per condition, spreading the generated embedding distribution and reducing retrieval accuracy. In the  
 333 context of AV simulation, diversity within valid behaviors is a feature — a single ego trajectory should  
 334 plausibly lead to multiple pedestrian behaviors (crossing, waiting, accelerating). R@1 measures only  
 335 conditioning alignment, not generation quality; FID captures the latter more faithfully.

### 336 Limitations.

- 337 • **Body pose quality:** Pedestrian poses are estimated (not ground-truth 3D), introducing noise  
 338 in training labels. Direct 3D body annotation from datasets like JRDB-Pose or BEDLAM  
 339 would improve quality.




Figure 4: H4 training trajectory: training-time validation FID (left axis, red) and R@1 (right axis, blue) across the late training segment. Epoch 3399 is the best checkpoint (highest R@1, FID within noise of epoch 3199, and best definitive 3-replication FID); both metrics then degrade steadily through epoch 4999. Training loss is flat across this range — sampling quality oscillates independently of training loss.

- **Conditioning scope:** We condition only on 2D ego trajectory. Incorporating relative pedestrian-ego geometry, ego velocity profile, or map elements could improve fidelity.
- **R@1 gap:** H4 achieves R@1 = 0.548, below H2’s 0.671. The adapter-style approach described above is a natural next step.
- **Physics:** Generated motions are not physically constrained (foot sliding, penetration). Post-hoc physics cleanup could improve realism for simulation deployment.
- **Evaluator domain shift:** FID and Diversity rely on motion-feature encoders pretrained on HumanML3D, whose motion distribution differs from in-the-wild pedestrian motion; absolute FID values should be compared only within our benchmark.
- **Single benchmark, no perceptual study:** All results are on our combined AVA/nuScenes/Waymo benchmark, and we do not yet report a human perceptual study of motion realism.

**Broader impact.** The primary intended use of ego-conditioned pedestrian motion generation is autonomous-driving simulation: synthesizing realistic, ego-aware pedestrian behaviors enables larger-scale safety testing without exposing real pedestrians to risk, and is particularly valuable for the long tail of rare or safety-critical interactions that are unsafe or impractical to collect. We note two risks. First, generated pedestrian motions could in principle be used adversarially, e.g. to seed motion-planning policies that learn to expect only nominal behaviors and consequently fail on out-of-distribution pedestrians; we mitigate this by reporting MultiModality and diversity metrics, which expose the breadth of the generated distribution rather than only its mean. Second, training data inherit demographic biases of the source AV datasets (AVA, nuScenes, Waymo): generated motion is conditioned on the distribution of pedestrians actually observed in U.S./European urban environments, so deployment in other settings should be preceded by additional validation. Generated motion data is synthetic and contains no personally identifiable information.

## 7 Conclusion

We introduced EgoPed, an ego-conditioned pedestrian body motion generation system based on latent diffusion with temporal cross-attention. Our key finding is that attending to the full 196-token ego trajectory sequence via cross-attention dramatically outperforms single-token pooled conditioning: 49% FID improvement (6.603 → 3.392) with non-overlapping confidence intervals. A paired ablation

reveals that frozen ego encoders are critical: unfreezing causes 34% MultiModality collapse as the encoder over-specializes to discriminative features.

These results establish a clear design recipe for ego-conditioned motion generation in AV settings: train the ego encoder contrastively on frozen VAE latents, then freeze it during diffusion training and use a cross-attention decoder as the denoiser. This recipe produces state-of-the-art ego-conditioned full-body pedestrian motions across AVA, nuScenes, and Waymo, and provides a foundation for scalable, behavior-consistent pedestrian simulation.

## References

- Xin Chen, Biao Jiang, Wen Liu, Zilong Huang, Bin Fu, Tao Chen, and Gang Yu. Executing your commands via motion diffusion in latent space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2212.04048.
- Ziyu Chen, Jiawei Yang, Jiahui Huang, Riccardo de Lutio, Janick Martinez Esturo, Boris Ivanovic, Or Litany, Zan Gojcic, Sanja Fidler, Marco Pavone, Li Song, and Yue Wang. Omnire: Omni urban scene reconstruction. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Lan Feng, Mohammadhossein Bahari, Kaouther Messaoud Ben Amor, Éloi Zablocki, Matthieu Cord, and Alexandre Alahi. UniTraj: A unified framework for scalable vehicle trajectory prediction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2024. arXiv:2403.15098.
- Chuan Guo, Shihao Zou, Xinxin Zuo, Sen Wang, Wei Ji, Xingyu Li, and Li Cheng. Generating diverse and natural 3d human motions from text. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Chuan Guo, Yuxuan Mu, Muhammad Gohar Javed, Sen Wang, and Li Cheng. MoMask: Generative masked modeling of 3d human motions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2312.00063.
- Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- Han Liang, Wenqian Zhang, Wenxuan Li, Jingyi Yu, and Lan Xu. InterGen: Diffusion-based multi-human motion generation under complex interactions. *International Journal of Computer Vision*, 2024. arXiv:2304.05684.
- Matthew Loper, Naureen Mahmood, Javier Romero, Gerard Pons-Moll, and Michael J. Black. SMPL: A skinned multi-person linear model. *ACM Transactions on Graphics, (Proc. SIGGRAPH Asia)*, 34(6):248:1–248:16, October 2015.
- Weibo Mao, Chenxin Xu, Qi Zhu, Siheng Chen, and Yanfeng Wang. Leapfrog diffusion model for stochastic trajectory prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2303.10895.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. arXiv:2112.10752.
- Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. DreamBooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2208.12242.
- Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. In *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2010.02502.
- Guy Tevet, Sigal Raab, Brian Gordon, Yonatan Shafir, Daniel Cohen-Or, and Amit H. Bermano. Human motion diffusion model. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2209.14916.

- 416 Darren Tsai, Julie Stephany Berrio, Mao Shan, Eduardo Nebot, and Stewart Worrall. Ms3d: Leverag-  
 417 ing multiple detectors for unsupervised domain adaptation in 3d object detection. In *2023 IEEE*  
 418 *26th International Conference on Intelligent Transportation Systems (ITSC)*, pages 140–147, 2023.  
 419 doi: 10.1109/ITSC57777.2023.10421941.
- 420 Darren Tsai, Julie Stephany Berrio, Mao Shan, Eduardo Nebot, and Stewart Worrall. Ms3d++:  
 421 Ensemble of experts for multi-source unsupervised domain adaptation in 3d object detection. *IEEE*  
 422 *Transactions on Intelligent Vehicles*, pages 1–16, 2024. doi: 10.1109/TIV.2024.3441527.
- 423 Hu Ye, Jun Zhang, Sibio Liu, Xiao Han, and Wei Yang. IP-Adapter: Text compatible image prompt  
 424 adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.
- 425 Jianrong Zhang, Yangsong Zhang, Xiaodong Cun, Yong Zhang, Hao Zhao, Hongtao Lu, Xi Shen,  
 426 and Ying Shan. T2M-GPT: Generating human motion from textual descriptions with discrete  
 427 representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*  
 428 *Recognition (CVPR)*, 2023a. arXiv:2301.06052.
- 429 Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image  
 430 diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*  
 431 *(ICCV)*, 2023b. arXiv:2302.05543.
- 432 Mingyuan Zhang, Xinyi Guo, Liang Pan, Zhongang Cai, Fangzhou Hong, Huirong Li, Lei Yang, and  
 433 Ziwei Liu. ReMoDiffuse: Retrieval-augmented motion diffusion model. In *Proceedings of the*  
 434 *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023c. arXiv:2304.01116.
- 435 Mingyuan Zhang, Zhongang Cai, Liang Pan, Fangzhou Hong, Xinyi Guo, Lei Yang, and Ziwei Liu.  
 436 MotionDiffuse: Text-driven human motion generation with diffusion model. *IEEE Transactions*  
 437 *on Pattern Analysis and Machine Intelligence*, 2024. arXiv:2208.15001.

## 438 A Implementation Details

439 **Ego encoder.** 4-layer transformer encoder, 4 attention heads, 256D hidden dimension (1024D feed-  
 440 forward), sinusoidal positional encoding for  $T = 196$  timesteps. Input: 2D ego position sequence,  
 441 projected linearly to 256D. No projection head (we verified the 2-layer MLP projection head does not  
 442 improve val MSE: 1.919 vs. 1.920). Pretraining: AdamW ( $\text{lr} = 1 \times 10^{-4}$ , weight decay  $10^{-4}$ ), 200  
 443 epochs, batch size 64, linear warmup followed by cosine annealing, gradient clipping at max-norm  
 444 1.0.

445 **Diffusion model.** MLD denoiser, `trans_dec` architecture: 9 layers, 4 attention heads, 256D token  
 446 dimension, 1024D feed-forward, latent sequence length  $L = 4$ . Diffusion: 1000 DDPM training steps  
 447 (scaled-linear  $\beta$  schedule), DDIM inference (50 steps,  $\eta = 0$ ). CFG dropout rate: 10%. Optimizer:  
 448 AdamW,  $\text{lr} = 1 \times 10^{-4}$  (constant), 5000 total epochs. Batch size: 128. Training hardware: NVIDIA  
 449 H100 (80 GB) on an academic HPC cluster; training runs in segments of up to 24h wall time.

450 **Checkpoint selection.** We select checkpoints based on validation FID (not training loss). For  
 451 H4, epoch 3399 is best; we confirm this with the full training trajectory and definitive 3-replication  
 452 evaluation. For H2, epoch 4399 is best despite continued training to epoch 5000.

453 **Evaluation.** All evaluations use `TEST.REPLICATION_TIMES=3` (3 independent generation passes),  
 454 `DIVERSITY_TIMES=300`, `MM_NUM_SAMPLES=100`, `MM_NUM_TIMES=10`. Each replication regenerates  
 455 the full test set with fresh sampling noise; the diffusion sampler is stochastic across replications while  
 456 the data pipeline and evaluator are fixed (training uses seed 1234). Confidence intervals are reported  
 457 as  $\pm$  CI (half-width of the 95% interval across replications).

## 458 B H6 Training Trajectory Table

## 459 C NeurIPS Paper Checklist

- 460 1. **Claims.** *Do the main claims made in the abstract and introduction accurately reflect the*  
 461 *paper’s contributions and scope?*

Table 6: H6 (FREEZE\_EGO=False) training-time validation metrics at selected checkpoint evaluation points (validation runs every 100 epochs; full curve in Figure 3). Training-time FID is best at epoch 2099 (4.789); R@1 rises until epoch  $\approx 3299$ , then saturates in the 0.37–0.42 band (maximum 0.417 at epoch 4599) while FID remains between 5.1 and 6.3.

| Epoch       | Train-time FID ↓ | Train-time R@1 ↑ |
|-------------|------------------|------------------|
| 99          | 9.835            | 0.036            |
| 499         | 8.373            | 0.038            |
| 999         | 6.011            | 0.065            |
| 1499        | 5.053            | 0.158            |
| 1999        | 4.967            | 0.271            |
| <b>2099</b> | <b>4.789</b>     | 0.273            |
| 2499        | 5.385            | 0.332            |
| 2999        | 5.889            | 0.373            |
| 3299        | 5.273            | 0.398            |
| 3399        | 5.141            | 0.380            |
| 3499        | 5.415            | 0.367            |
| 3999        | 6.221            | 0.407            |
| <b>4599</b> | 6.339            | <b>0.417</b>     |
| 4999        | 5.561            | 0.408            |

- 462 Answer: [Yes]. Justification: The two headline claims (49% FID improvement from  
463 temporal cross-attention; MultiModality collapse from unfreezing the ego encoder) are  
464 supported by replicated evaluations with confidence intervals in Tables 3, 4, and 5.
- 465 2. **Limitations.** *Does the paper discuss the limitations of the work?*  
466 Answer: [Yes]. Justification: Section 6 lists estimated (not ground-truth) pose labels,  
467 restricted conditioning scope, the R@1 gap, lack of physical constraints, evaluator domain  
468 shift, and the absence of a perceptual study.
- 469 3. **Theory.** *Does the paper include theoretical results?*  
470 Answer: [NA]. Justification: This is an empirical systems paper; no theoretical claims are  
471 made.
- 472 4. **Experimental reproducibility.** *Does the paper fully disclose the information needed to*  
473 *reproduce the main experimental results?*  
474 Answer: [Yes]. Justification: Architecture, training, and evaluation hyperparameters are  
475 given in Section 4 and Appendix A; the dataset construction pipeline is described in Sec-  
476 tion 4.1.
- 477 5. **Code and data access.** *Does the paper provide open access to data and code?*  
478 Answer: [No]. Justification: Code and the derived dataset will be released upon publication;  
479 the underlying nuScenes and Waymo data are publicly available under their research licenses,  
480 and the AVA recordings will be released subject to anonymization requirements.
- 481 6. **Experimental settings.** *Does the paper specify all training and test details?*  
482 Answer: [Yes]. Justification: Optimizer, learning rate, batch size, epochs, CFG dropout,  
483 sampler settings, and checkpoint-selection protocol are in Appendix A; evaluation settings  
484 (replications, metric parameters) in Section 4.2 and Appendix A.
- 485 7. **Error bars.** *Does the paper report error bars or other statistical significance information?*  
486 Answer: [Yes]. Justification: All key metrics report 95% confidence intervals over 3  
487 independent generation replications; the headline FID comparison has non-overlapping  
488 intervals.
- 489 8. **Compute resources.** *Does the paper provide sufficient information on the compute re-*  
490 *sources?*  
491 Answer: [Yes]. Justification: Single NVIDIA H100 (80 GB) per run,  $\leq 24$ h wall-time  
492 segments,  $\sim 2$  segments per diffusion model (Appendix A); evaluations take  $\sim 8$  minutes  
493 each.
- 494 9. **Code of ethics.** *Does the research conform to the NeurIPS Code of Ethics?*  
495 Answer: [Yes]. Justification: The work uses licensed public AV datasets and our own staged

496 recordings with consenting participants; generated data contains no personally identifiable  
497 information.

498 10. **Broader impacts.** *Does the paper discuss potential positive and negative societal impacts?*  
499 Answer: [Yes]. Justification: Section 6 discusses safety benefits for AV simulation,  
500 adversarial-use and dataset-bias risks, and mitigations.

501 11. **Safeguards.** *Does the paper describe safeguards for responsible release?*  
502 Answer: [NA]. Justification: The model generates pedestrian body motion for simulation  
503 and poses no high misuse risk requiring gated release.

504 12. **Licenses.** *Are existing assets properly credited and licenses respected?*  
505 Answer: [Yes]. Justification: nuScenes and Waymo are used under their research licenses  
506 and cited; the MLD backbone, HumanML3D representation and evaluators, OmniRe, MS3D,  
507 and SMPL are cited.

508 13. **New assets.** *Are new assets well documented?*  
509 Answer: [Yes]. Justification: The dataset construction pipeline (annotation, pose extraction,  
510 interaction cropping) is documented in Section 4.1; documentation will accompany the  
511 release.

512 14. **Human subjects.** *Does the paper involve crowdsourcing or research with human subjects?*  
513 Answer: [Yes]. Justification: Staged interaction scenarios were recorded with informed,  
514 consenting participants; recordings follow institutional data-protection procedures, and  
515 released data will be anonymized.