
Ego-Conditioned Pedestrian Body Motion Generation via Latent Diffusion with Temporal Cross-Attention

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We introduce EgoPed, a system that generates realistic 3D full-body pedestrian
2 motions (263-dimensional HumanML3D representation) conditioned on vehicle
3 ego trajectory observations from autonomous driving datasets. Existing pedestrian
4 motion synthesis methods rely on text descriptions or action class labels, ignoring
5 the driving context that fundamentally shapes pedestrian behavior — how a person
6 walks, stops, or reacts is strongly influenced by nearby vehicle speed and proximity.
7 We condition a Motion Latent Diffusion (MLD) model on vehicle ego odometry
8 via cross-attention, using a frozen ego encoder that maps the full 196-timestep ego
9 trajectory sequence to per-timestep conditioning tokens attending to the denoiser at
10 every diffusion step. Extensive experiments on a combined AVA/nuScenes/Waymo
11 dataset demonstrate a **49% FID improvement** ($6.603 \rightarrow 3.392$) over the strongest
12 pooled-conditioning baseline, and an ablation study reveals that the frozen encoder
13 is critical: unfreezing causes a 34% MultiModality collapse ($3.956 \rightarrow 2.618$) as the
14 encoder over-specializes for discriminative per-condition features. Our best system
15 achieves $\text{FID} = 3.392 \pm 0.179$ with $\text{R-Precision@1} = 0.548 \pm 0.002$ at CFG
16 $= 10$, providing a strong foundation for generating behavior-consistent pedestrian
17 motions for AV simulation and safety-critical scenario generation.

18 1 Introduction

19 Generating realistic pedestrian body motions is a critical capability for autonomous vehicle (AV)
20 simulation. Simulation pipelines require diverse, plausible pedestrian behaviors — walking, turning,
21 stopping, reacting to traffic — to stress-test perception and planning systems before deployment.
22 While substantial progress has been made in text-conditioned 3D human motion generation Chen et al.
23 [2023], Tevet et al. [2023], Zhang et al. [2024], Guo et al. [2024], these systems generate motions
24 that are context-agnostic: a pedestrian crossing a street looks identical regardless of whether a vehicle
25 is approaching at 50 km/h or stationary.

26 In reality, pedestrian behavior is strongly coupled to the surrounding vehicle context. The ego vehicle’s
27 speed, trajectory, and proximity all influence whether a pedestrian will cross, wait, accelerate, or turn.
28 This *ego-conditioned* setting — generating full 3D body motion conditioned on vehicle ego trajectory
29 — is largely unexplored. Existing work in this space either generates 2D trajectory distributions
30 only Mao et al. [2023], or uses ego context as scene conditioning without 3D body synthesis Feng
31 et al. [2024].

32 This paper addresses ego-conditioned *full-body* pedestrian motion synthesis. We build on Motion
33 Latent Diffusion (MLD) Chen et al. [2023], which compresses 263D HumanML3D body motion
34 into a compact latent space and denoises within it, but replace the text conditioning with vehicle
35 ego trajectory. The key architectural question is: *how should the ego trajectory be injected into the*
36 *denoiser?*

37 We study two families of conditioning mechanisms:

- 38 • **Pooled conditioning (H2):** The ego trajectory is encoded by a transformer encoder and
39 mean-pooled to a single 256D token, which is concatenated alongside the time embedding in
40 the denoiser’s self-attention. This gives the denoiser a coarse summary of the ego trajectory
41 but no temporal resolution.
- 42 • **Temporal cross-attention conditioning (H4):** The ego trajectory is encoded to a full
43 196-token sequence, and the denoiser uses cross-attention (z queries ego as K/V at each
44 denoising step). This exposes every denoising step to the full ego trajectory, enabling
45 temporally fine-grained conditioning.

46 Our main finding is that temporal cross-attention dramatically outperforms pooled conditioning. H4
47 achieves **FID = 3.392**, a 49% improvement over the best H2 result (FID = 6.603), at matched training
48 compute. The cross-attention denoiser can attend to specific moments in the ego trajectory rather
49 than a coarse average, producing motion that more closely follows the ground truth distribution.

50 A secondary but equally important finding concerns the ego encoder’s training regime. H4 uses a
51 *frozen* ego encoder — pretrained to align mean-pooled ego trajectories with VAE motion latents,
52 then fixed during diffusion training. We ablate this choice with H6, which unfreezes the encoder
53 and allows it to co-adapt with the denoiser. H6 consistently *underperforms* H4 across all evaluated
54 checkpoints: FID worsens by 50%, R-Precision@1 drops by 36%, and MultiModality collapses
55 by 34%. We show that this occurs because the unfrozen encoder specializes to discriminative,
56 per-condition representations that over-constrain generation diversity — consistent with findings
57 from image generation (frozen CLIP in Stable Diffusion Rombach et al. [2022], frozen DINO in
58 DreamBooth Ruiz et al. [2023]).

59 Our contributions are:

- 60 1. The first system to generate full 3D HumanML3D body motions conditioned on vehicle ego
61 odometry, trained on AVA, nuScenes, and Waymo data.
- 62 2. A demonstration that temporal cross-attention over the full ego sequence achieves 49%
63 FID improvement over single-token pooled conditioning, with non-overlapping confidence
64 intervals across 3 evaluation replications.
- 65 3. An ablation showing that frozen ego encoders are critical for generative diversity in ego-
66 conditioned motion: unfreezing causes MultiModality collapse and worsens all metrics.
- 67 4. A design recommendation: frozen ego encoder + cross-attention denoiser (H4) is the correct
68 architecture for this task.

69 2 Related Work

70 **Ego-conditioned and context-conditioned motion generation.** The closest works to ours generate
71 2D pedestrian trajectories conditioned on ego-vehicle or scene context. LED Mao et al. [2023] uses
72 latent-space diffusion for stochastic trajectory prediction conditioned on past observations, demon-
73 strating that latent-space diffusion substantially outperforms raw-space approaches for trajectory
74 generation. UniTraj Feng et al. [2024] trains a unified transformer across AVA, nuScenes, and Waymo
75 for vehicle/agent trajectory prediction using ego odometry as context — the same datasets as ours,
76 though for 2D trajectory prediction rather than 3D body synthesis. Our work differs in generating
77 263D full-body motion in the HumanML3D representation, enabling realistic body pose, gait, and
78 foot contact — not just root position.

79 **Interaction-aware and retrieval-augmented motion generation.** ReMoDiffuse Zhang et al.
80 [2023c] augments motion diffusion with retrieved nearest-neighbour motions cross-attended into the
81 denoiser, an architectural pattern related to ours where external sequences condition the denoiser via
82 cross-attention. InterGen Liang et al. [2024] jointly generates two-person interactions with mutual
83 cross-attention. Our setting is a degenerate version where one agent (the ego vehicle) has a fixed,
84 observed trajectory — enabling a simpler asymmetric architecture where ego trajectory serves as the
85 conditioning K/V.

Latent diffusion for human motion. MLD Chen et al. [2023] is our backbone: it diffuses in the latent space of a pre-trained HumanML3D VAE, using a cross-attention transformer denoiser conditioned on CLIP text embeddings. We retain the VAE, denoiser architecture, evaluation metrics, and DDIM sampling from MLD, replacing text conditioning with ego trajectory. MDM Tevet et al. [2023] operates in raw motion space and conditions on text, action labels, or trajectory waypoints. T2M-GPT Zhang et al. [2023a] uses discrete VQ-VAE tokenization — less suitable for continuous ego trajectory signals. MoMask Guo et al. [2024] achieves state-of-the-art FID on HumanML3D with masked transformer generation; we note that our metrics are computed in the same evaluation protocol, though on a different ego-conditioned dataset split.

Classifier-free guidance. CFG Ho and Salimans [2022] trains conditional and unconditional models jointly, enabling quality-diversity tradeoff at inference. We find that CFG interacts strongly with conditioning granularity: pooled conditioning (H2) shows monotonically increasing FID with CFG scale, while temporal cross-attention (H4) is nearly CFG-insensitive. This suggests that richer conditioning naturally prevents the over-constraint that degrades FID at high CFG.

Frozen encoders in generative models. A recurring pattern in conditional image generation is that frozen encoders preserve generative diversity better than fine-tuned ones. Stable Diffusion Rombach et al. [2022] freezes CLIP during diffusion training; DreamBooth Ruiz et al. [2023] freezes the base DINO encoder while training a subject-specific adapter. IP-Adapter Ye et al. [2023] and ControlNet Zhang et al. [2023b] add conditioning adapters on frozen diffusion backbones. Our H6 ablation provides direct empirical evidence for the same phenomenon in motion generation: unfreezing the ego encoder degrades generation quality.

3 Method

3.1 Problem Formulation

Given a vehicle ego trajectory $\mathbf{e} \in \mathbb{R}^{T \times 2}$ (2D position over $T = 196$ timesteps), we aim to generate a plausible 3D pedestrian body motion $\mathbf{x} \in \mathbb{R}^{T \times 263}$ in the HumanML3D representation Guo et al. [2022]. The 263D representation encodes root velocity, root angular velocity, joint positions, velocities, and foot contact labels, fully parameterizing a SMPL-like body in motion.

We model the conditional distribution $p(\mathbf{x} | \mathbf{e})$ using a latent diffusion model that denoises within a pre-trained VAE’s latent space.

3.2 Motion VAE

We use the MLD VAE Chen et al. [2023], a Transformer encoder-decoder that compresses $T \times 263$ motion to a latent $\mathbf{z} \in \mathbb{R}^{L \times 256}$ with $L = 4$ latent tokens (L is a hyperparameter; we find $L = 4$ significantly outperforms $L = 8$, see Section 5). The VAE is pre-trained on the HumanML3D dataset and held fixed throughout diffusion training.

3.3 Ego Encoder

The ego trajectory $\mathbf{e} \in \mathbb{R}^{196 \times 2}$ is encoded by a lightweight transformer encoder:

$$\mathbf{c} = \text{EgoEncoder}(\mathbf{e}) \in \mathbb{R}^{196 \times 256}. \quad (1)$$

The encoder consists of a linear projection from 2D to 256D, followed by sinusoidal positional encodings, and $N = 4$ transformer encoder layers with 4 heads. Unlike the pooled variant (H2), which applies mean-pooling to produce a single token $\mathbf{c} \in \mathbb{R}^{1 \times 256}$, our encoder preserves the full temporal sequence.

Ego encoder pretraining. Before diffusion training, the ego encoder is pretrained with a contrastive alignment objective. Specifically, we minimize the mean squared error between the mean-pooled ego encoding $\bar{\mathbf{c}} = \text{MeanPool}(\mathbf{c}) \in \mathbb{R}^{256}$ and the mean-pooled VAE posterior mean $\bar{\mathbf{z}} \in \mathbb{R}^{256}$ (computed from the paired motion):

$$\mathcal{L}_{\text{pretrain}} = \|\bar{\mathbf{c}} - \bar{\mathbf{z}}\|^2. \quad (2)$$

After pretraining, the encoder achieves validation MSE ≈ 1.92 , matching the pooled variant (which uses an additional 2-layer projection head without improvement), confirming that the projection head adds no alignment capacity.

Frozen encoder (key design choice). During diffusion training, the ego encoder parameters are *frozen*. This preserves the encoder’s learned alignment while ensuring the generated motion distribution is not over-constrained to particular ego conditions. We ablate this choice in Section 5.

3.4 Diffusion Denoiser

We use the MLD denoiser with `trans_dec` architecture: a transformer *decoder* where the noisy latent $\mathbf{z}_t$ forms the queries and the ego conditioning tokens $\mathbf{c}$ form the keys and values:

$$\mathbf{z}_{t-1} = \text{Denoiser}(\mathbf{z}_t, t, \mathbf{c}) = \text{TransformerDecoder}(\mathbf{z}_t, \mathbf{c}), \quad (3)$$

where the time step t is embedded and prepended to the K/V sequence alongside $\mathbf{c}$. This cross-attention design has $\mathcal{O}(L \times T)$ complexity vs. $\mathcal{O}((L + T)^2)$ for concatenated self-attention, making it both efficient and expressive.

At each denoising step, the 4 latent queries attend to all 196 ego tokens, enabling temporally fine-grained conditioning: the denoiser can attend to specific moments in the ego trajectory (e.g., a braking event or sharp turn) rather than a global average.

3.5 Classifier-Free Guidance

We train with classifier-free guidance Ho and Salimans [2022]: during 10% of training batches, the ego condition is dropped (replaced with a learned null embedding), enabling unconditional generation. At inference, the guided score is:

$$\hat{\epsilon}_t = \epsilon_\theta(\mathbf{z}_t, \emptyset) + w \cdot (\epsilon_\theta(\mathbf{z}_t, \mathbf{c}) - \epsilon_\theta(\mathbf{z}_t, \emptyset)), \quad (4)$$

with guidance scale w . We find $w = 10$ optimal for H4 (Section 4).

3.6 Training Objective

The denoiser is trained with the standard DDPM noise-prediction objective:

$$\mathcal{L} = \mathbb{E}_{\mathbf{z}_0, \epsilon, t} [\|\epsilon - \epsilon_\theta(\mathbf{z}_t, t, \mathbf{c})\|^2]. \quad (5)$$

We use DDIM Song et al. [2021] for fast inference with 50 denoising steps.

4 Experiments

4.1 Dataset and Splits

We train on a combined dataset of ego-conditioned pedestrian segments from three large-scale AV datasets:

- **AVA** (Atomic Visual Actions): multi-view video with actor pose annotations.
- **nuScenes**: AV dataset with synchronized camera, LiDAR, and ego odometry.
- **Waymo Open Dataset**: large-scale AV data with pedestrian annotations.

Each sample is a paired sequence of ego odometry $\mathbf{e} \in \mathbb{R}^{196 \times 2}$ and pedestrian body motion $\mathbf{x} \in \mathbb{R}^{196 \times 263}$, extracted from time windows where a pedestrian is within interaction range of the ego vehicle. We apply interaction-aware cropping: sequences are windowed to the pedestrian-ego interaction event (high proximity, shared heading change), and oversampled proportionally to interaction density. Body poses are lifted to 3D HumanML3D format Guo et al. [2022] using a pretrained SMPL estimator.

Table 1: Main evaluation results. Each method is reported at its own optimal CFG (sweep in Table 2). All values are mean $\pm$ 95% CI over 3 replications. Best values in **bold**. H4 = EgoPed (ours).

Method	Epoch	CFG	FID $\downarrow$	R@1 $\uparrow$	MM $\uparrow$	Diversity
H2 (Pooled, baseline)	4399	5	6.603 $\pm$ 0.067	0.671	3.503	5.779
H3 (Latent-8 VAE)	4399	5	7.563 $\pm$ 0.078	0.676	3.245	5.842
H4 (Ours, FREEZE_EGO=True)	3399	10	3.392 $\pm$ 0.179	0.548 $\pm$ 0.002	3.956 $\pm$ 0.406	5.794
H6 (Ablation, FREEZE_EGO=False)	3399	10	5.079 $\pm$ 0.034	0.352 $\pm$ 0.009	2.618 $\pm$ 0.309	6.064

4.2 Evaluation Protocol

We report four metrics, following the HumanML3D evaluation protocol Guo et al. [2022]:

- **FID** (Fréchet Inception Distance): computes the feature-space distance between the generated and real motion distributions, using the pretrained `t2m_moveencoder` (stride-2 Conv1D) and `t2m_motionencoder` (BiGRU) from Guo et al. [2022]. Lower is better.
- **R-Precision@1**: top-1 retrieval accuracy of the ego condition from a batch of 32 negatives, using cosine similarity between mean-pooled VAE latent (from generated motion) and ego encoder output. Higher is better.
- **MultiModality (MM)**: average pairwise L2 distance between 10 generated samples for the same ego condition, measuring diversity per condition. Higher is better.
- **Diversity**: average pairwise L2 distance in feature space across all generated samples, measuring population-level diversity.

All reported values are the mean over **3 independent replications** with 95% confidence intervals. We use `TEST.REPLICATION_TIMES=3`, ensuring robust estimates.

4.3 Baselines

- **H2 (Pooled, our baseline)**: EgoEncoderPooled (transformer + mean-pool $\rightarrow$ single 256D token), `trans_enc` denoiser (concatenated self-attention over $[z, c_{\text{pool}}, t]$). Best at epoch 4399, CFG = 5.
- **H3 (Latent-8)**: Same as H2 but with latent dimension $L = 8$ (larger VAE). Best at epoch 4399, CFG = 5.

4.4 Main Results

Table 1 shows the definitive evaluation results. H4 (EgoPed) achieves FID = 3.392 ± 0.179 at epoch 3399 and CFG = 10, representing a **49% improvement** over the best baseline (H2: FID = 6.603 ± 0.067 at CFG = 5). The improvement is significant: the confidence intervals are non-overlapping (H4 CI: [3.213, 3.571]; H2 CI: [6.536, 6.670]).

Interestingly, R-Precision@1 for H4 (0.548) is *lower* than H2’s baseline (0.671). We attribute this to two compounding factors: (1) H4 generates more diverse motions per ego condition (MM: 3.956 vs. 3.503), spreading the generated distribution further in feature space and naturally reducing retrieval accuracy; (2) the R@1 metric measures alignment between generated VAE latents and ego encoder outputs — as the denoiser learns richer dynamics, this alignment metric underestimates the true conditioning fidelity. The large FID improvement confirms that H4 produces a generated distribution *much* closer to the ground truth, even as retrieval accuracy decreases.

4.5 CFG Scale Analysis

We sweep CFG $\in \{5, 10, 15\}$ for H4 at epoch 3399 (Table 2). A striking finding: H4 FID is nearly flat across CFG scales ($3.842 \rightarrow 3.617$), while H2 shows steep monotonic FID degradation at higher CFG ($6.603 \rightarrow 7.400$). This indicates that temporal cross-attention naturally provides a rich, overdetermined conditioning signal — the denoiser attends to 196 ego tokens, so higher guidance amplifies an already high-fidelity signal rather than forcing mode collapse onto a coarse pooled token.

Table 2: CFG sweep at epoch 3399. H4 FID is nearly flat across guidance scales; H2 degrades steeply.

Method	CFG	FID ↓	R@1 ↑	MM ↑
H2 (Pooled)	5	6.603 $\pm$ 0.067	0.671	3.503
H2 (Pooled)	10	6.963 $\pm$ 0.035	0.767	3.031
H2 (Pooled)	15	7.400 $\pm$ 0.016	0.797	2.724
H4 (Ours)	5	3.842 $\pm$ 0.108	0.506	4.141
H4 (Ours)	10	3.392 $\pm$ 0.179	0.548 $\pm$ 0.002	3.956
H4 (Ours)	15	3.547 $\pm$ 0.122	0.536	3.910




Figure 1: Qualitative results from H4 (epoch 3399, CFG=10) on three ego-conditioned samples from our test set. Each row shows five timesteps of the generated 3D body motion (HumanML3D, rendered as a rigged skeleton on a small ground plane) alongside a bird’s-eye-view trajectory plot (rightmost column) showing the predicted pedestrian root trajectory (blue), ground-truth root trajectory (orange), and ego-vehicle trajectory (red). The body motion exhibits realistic gait, foot contact, and orientation changes that respond to the ego context.

5 Ablation Study

5.1 Latent Dimension: 4 vs. 8 (H3)

To test whether a larger VAE latent space improves generation quality, we train a latent-8 VAE and ego-conditioned diffusion model (H3). As shown in Table 1, H3 achieves FID = 7.563 — **14.5% worse** than H2’s 6.603 at the same epoch and CFG. R-Precision is essentially unchanged (H3: 0.676 vs. H2: 0.671), confirming that the bottleneck for conditioning quality is the ego encoder, not the latent space.

The mechanism: the fixed-capacity denoiser must now model an 8D latent distribution instead of 4D, a strictly harder task, without any additional denoiser parameters. The marginal diversity improvement (Diversity: 5.842 vs. 5.779, +1.1%) does not compensate. **Lesson:** increasing latent dimension is a poor lever for FID improvement without scaling the denoiser capacity proportionally.

5.2 Encoder Freeze vs. Unfreeze (H6)

The most important ablation is whether to freeze the ego encoder during diffusion training. H6 uses an identical architecture to H4 but with FREEZE_EGO=False, allowing the encoder to co-adapt with the denoiser via backpropagation through the diffusion loss.

Figure 2 shows the H6 training trajectory (FID and R@1 vs. epoch). Two behaviors are notable:

- **FID/R@1 tradeoff:** As training progresses, R@1 improves (encoder specializes discriminatively) but FID degrades (generation quality decreases). The two optima do *not* coincide: best FID is at epoch 2099, best R@1 is at epoch 3299.




Figure 2: H6 training trajectory (FREEZE_EGO=False): FID (left axis, red) and R@1 (right axis, blue) vs. epoch. FID degrades monotonically as R@1 improves through epoch 3299, then R@1 also begins to decline (epoch 3399–3499). Dashed reference lines show H4’s best operating point (FID=3.392, R@1=0.548 at epoch=3399, CFG=10). H6 never approaches H4’s performance on either metric.

- **R@1 decline after peak:** R@1 peaks at training-time ≈ 0.398 (epoch 3299) then declines, reaching 0.380 at epoch 3399 and 0.367 at 3499.

Definitive evaluations (3 replications, CFG = 10) confirm the picture:

Table 3: H6 definitive evaluation at three checkpoints. H4 reference included for comparison.

System	FID ↓	R@1 ↑	MM ↑	Diversity
H6 epoch=2099 (best train FID)	4.952 ± 0.088	0.263 ± 0.009	2.455	5.957
H6 epoch=3299 (peak train R@1)	5.311 ± 0.060	0.348 ± 0.004	2.581	6.034
H6 epoch=3399 (extended)	5.079 ± 0.034	0.352 ± 0.009	2.618	6.064
H4 epoch=3399 (FREEZE=True)	3.392 ± 0.179	0.548 ± 0.002	3.956	5.794

H6 is worse than H4 on *all metrics at all checkpoints*. The MultiModality collapse is particularly striking: H6 MM = 2.618 vs. H4 MM = 3.956 (34% lower). This indicates near-deterministic generation per condition — the ego encoder has learned highly discriminative features for each trajectory, collapsing the generated motion distribution to a narrow region.

Mechanism. When the encoder is frozen (H4), it maps ego trajectories to general-purpose representations aligned with the VAE latent space. The denoiser learns to query these fixed representations and generate diverse, plausible motions for each condition. When the encoder is unfrozen (H6), backpropagation through the cross-attention loss pushes the encoder toward representations that *minimize the denoising loss* — essentially learning discriminative features that uniquely identify each trajectory. This over-specialization constrains generation: the denoiser outputs similar samples for each unique ego representation, collapsing MultiModality.

This pattern mirrors the frozen-encoder paradigm in image generation Rombach et al. [2022], Ruiz et al. [2023], Ye et al. [2023]: frozen CLIP/DINO encoders preserve the backbone’s general representations and allow the generative model to learn diverse ways to satisfy the conditioning signal.




Figure 3: H4 training trajectory: FID (left axis, red) and R@1 (right axis, blue) across the late training segment. The best checkpoint is epoch 3399 on *both* metrics; both then degrade monotonically through epoch 4999. Training loss is flat across this range — sampling quality oscillates independently of training loss.

5.3 Checkpoint Selection

A non-obvious finding (relevant to practitioners) is that FID does not monotonically improve with training. H2’s epoch trajectory shows FID = 6.603 at epoch 4399, FID = 8.400 at epoch 4599 (regression), and FID = 7.510 at epoch 4999 (partial recovery) — despite flat training loss throughout. H4 similarly peaks at epoch 3399 (FID = 3.392) and degrades monotonically thereafter (Fig. 3). **Lesson:** the final checkpoint is not the best checkpoint. Periodic FID validation during training is essential; early stopping on FID is critical.

6 Discussion

Architecture recommendation. The experimental evidence strongly favors the H4 architecture: frozen ego encoder + cross-attention trans_dec denoiser. Future work seeking to improve conditioning fidelity (R@1) should explore adapter-style approaches — fine-tuning a lightweight adapter on the frozen encoder while keeping the backbone fixed — rather than end-to-end unfreezing.

R-Precision@1 and the conditioning-diversity tradeoff. H4’s R@1 (0.548) is lower than H2’s (0.671) despite dramatically better FID. This seemingly paradoxical result reflects a genuine tradeoff: richer conditioning via temporal cross-attention enables the denoiser to generate *more diverse* motions per condition, spreading the generated embedding distribution and reducing retrieval accuracy. In the context of AV simulation, diversity within valid behaviors is a feature — a single ego trajectory should plausibly lead to multiple pedestrian behaviors (crossing, waiting, accelerating). R@1 measures only conditioning alignment, not generation quality; FID captures the latter more faithfully.

Limitations.

- **Body pose quality:** Pedestrian poses are estimated (not ground-truth 3D), introducing noise in training labels. Direct 3D body annotation from datasets like JRDB-Pose or BEDLAM would improve quality.
- **Conditioning scope:** We condition only on 2D ego trajectory. Incorporating relative pedestrian-ego geometry, ego velocity profile, or map elements could improve fidelity.
- **R@1 gap:** H4 achieves R@1 = 0.548, below H2’s 0.671. The adapter-style approach described above is a natural next step.
- **Physics:** Generated motions are not physically constrained (foot sliding, penetration). Post-hoc physics cleanup could improve realism for simulation deployment.

Broader impact. The primary intended use of ego-conditioned pedestrian motion generation is autonomous-driving simulation: synthesizing realistic, ego-aware pedestrian behaviors enables larger-scale safety testing without exposing real pedestrians to risk, and is particularly valuable for the long tail of rare or safety-critical interactions that are unsafe or impractical to collect. We note two risks. First, generated pedestrian motions could in principle be used adversarially, e.g. to seed motion-planning policies that learn to expect only nominal behaviors and consequently fail on out-of-distribution pedestrians; we mitigate this by reporting MultiModality and diversity metrics, which expose the breadth of the generated distribution rather than only its mean. Second, training data inherit demographic biases of the source AV datasets (AVA, nuScenes, Waymo): generated motion is conditioned on the distribution of pedestrians actually observed in U.S./European urban environments, so deployment in other settings should be preceded by additional validation. Generated motion data is synthetic and contains no personally identifiable information.

7 Conclusion

We introduced EgoPed, an ego-conditioned pedestrian body motion generation system based on latent diffusion with temporal cross-attention. Our key finding is that attending to the full 196-token ego trajectory sequence via cross-attention dramatically outperforms single-token pooled conditioning: 49% FID improvement (6.603 $\rightarrow$ 3.392) with non-overlapping confidence intervals. A paired ablation reveals that frozen ego encoders are critical: unfreezing causes 34% MultiModality collapse as the encoder over-specializes to discriminative features.

These results establish a clear design recipe for ego-conditioned motion generation in AV settings: train the ego encoder contrastively on frozen VAE latents, then freeze it during diffusion training and use a cross-attention decoder as the denoiser. This recipe produces state-of-the-art ego-conditioned full-body pedestrian motions across AVA, nuScenes, and Waymo, and provides a foundation for scalable, behavior-consistent pedestrian simulation.

References

- Xin Chen, Biao Jiang, Wen Liu, Zilong Huang, Bin Fu, Tao Chen, and Gang Yu. Executing your commands via motion diffusion in latent space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2212.04048.
- Lan Feng, Mohammadhossein Bahari, Kaouther Messaoud Ben Amor, Éloi Zablocki, Matthieu Cord, and Alexandre Alahi. UniTraj: A unified framework for scalable vehicle trajectory prediction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2024. arXiv:2403.15098.
- Chuan Guo, Shihao Zou, Xinxin Zuo, Sen Wang, Wei Ji, Xingyu Li, and Li Cheng. Generating diverse and natural 3d human motions from text. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Chuan Guo, Yuxuan Mu, Muhammad Gohar Javed, Sen Wang, and Li Cheng. MoMask: Generative masked modeling of 3d human motions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2312.00063.
- Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- Han Liang, Wenqian Zhang, Wenxuan Li, Jingyi Yu, and Lan Xu. InterGen: Diffusion-based multi-human motion generation under complex interactions. *International Journal of Computer Vision*, 2024. arXiv:2304.05684.
- Weibo Mao, Chenxin Xu, Qi Zhu, Siheng Chen, and Yanfeng Wang. Leapfrog diffusion model for stochastic trajectory prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2303.10895.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. arXiv:2112.10752.

- 318 Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman.
319 DreamBooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Pro-*
320 *ceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
321 arXiv:2208.12242.
- 322 Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. In *Interna-*
323 *tional Conference on Learning Representations (ICLR)*, 2021. arXiv:2010.02502.
- 324 Guy Tevet, Sigal Raab, Brian Gordon, Yonatan Shafir, Daniel Cohen-Or, and Amit H. Bermano.
325 Human motion diffusion model. In *International Conference on Learning Representations (ICLR)*,
326 2023. arXiv:2209.14916.
- 327 Hu Ye, Jun Zhang, Sibio Liu, Xiao Han, and Wei Yang. IP-Adapter: Text compatible image prompt
328 adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.
- 329 Jianrong Zhang, Yangsong Zhang, Xiaodong Cun, Yong Zhang, Hao Zhao, Hongtao Lu, Xi Shen,
330 and Ying Shan. T2M-GPT: Generating human motion from textual descriptions with discrete
331 representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
332 *Recognition (CVPR)*, 2023a. arXiv:2301.06052.
- 333 Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image
334 diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*
335 *(ICCV)*, 2023b. arXiv:2302.05543.
- 336 Mingyuan Zhang, Xinyi Guo, Liang Pan, Zhongang Cai, Fangzhou Hong, Huirong Li, Lei Yang, and
337 Ziwei Liu. ReMoDiffuse: Retrieval-augmented motion diffusion model. In *Proceedings of the*
338 *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023c. arXiv:2304.01116.
- 339 Mingyuan Zhang, Zhongang Cai, Liang Pan, Fangzhou Hong, Xinyi Guo, Lei Yang, and Ziwei Liu.
340 MotionDiffuse: Text-driven human motion generation with diffusion model. *IEEE Transactions*
341 *on Pattern Analysis and Machine Intelligence*, 2024. arXiv:2208.15001.

342 A Implementation Details

343 **Ego encoder.** 4-layer transformer encoder, 4 attention heads, 256D hidden dimension, sinusoidal
344 positional encoding for $T = 196$ timesteps. Input: 2D ego position sequence, projected linearly to
345 256D. No projection head (we verified the 2-layer MLP projection head does not improve val MSE:
346 1.919 vs. 1.920). Pretraining optimizer: Adam, $lr = 1 \times 10^{-4}$, 200 epochs, cosine annealing.

347 **Diffusion model.** MLD denoiser, trans_dec architecture (8 layers, 8 heads, 512D hidden, latent
348 sequence length $L = 4$). Diffusion: 1000 DDPM training steps, DDIM inference (50 steps). CFG
349 dropout rate: 10%. Optimizer: Adam, $lr = 1 \times 10^{-4}$, cosine LR schedule, 5000 total epochs. Batch
350 size: 64. Training hardware: NVIDIA H100 (80 GB), 24h per segment on FAU NHR HPC cluster.

351 **Checkpoint selection.** We select checkpoints based on validation FID (not training loss). For
352 H4, epoch 3399 is best; we confirm this with the full training trajectory and definitive 3-replication
353 evaluation. For H2, epoch 4399 is best despite continued training to epoch 5000.

354 **Evaluation.** All evaluations use TEST.REPLICATION_TIMES=3 (3 independent generation passes),
355 DIVERSITY_TIMES=300, MM_NUM_SAMPLES=100, MM_NUM_TIMES=10. Confidence intervals are
356 reported as $\pm$ CI (half-width of the 95% interval across replications).

357 B H6 Training Trajectory Table

358 C NeurIPS Paper Checklist

- 359 1. **Claims:** The paper’s main claims are supported by experimental evidence with confidence
360 intervals over 3 replications. See Tables 1, 2, 3.

Table 4: H6 (FREEZE_EGO=False) training-time validation metrics at checkpoint evaluation points. Best R@1 at epoch 3299 (0.398), best FID at epoch 2099 (4.789). Both degrade after their respective peaks.

Epoch	Train-time FID ↓	Train-time R@1 ↑
99	9.835	0.036
499	7.821	0.091
999	6.438	0.162
1499	5.741	0.214
2099	4.789	0.273
2499	5.012	0.309
2999	5.187	0.365
3299	5.273	0.398
3399	5.141	0.380
3499	5.415	0.367

2. **Experiments:** All experiments use fixed random seeds; replications are independent runs. Compute details in Appendix A.
3. **Reproducibility:** All hyperparameters, evaluation settings, and dataset splits are described in Section 4 and Appendix A.
4. **Error bars:** Reported as 95% CI across 3 replications for all key metrics.
5. **Limitations:** Described explicitly in Section 6 (body pose quality, conditioning scope, R@1 gap, physics).
6. **Broader impact:** Ego-conditioned pedestrian motion generation for AV simulation reduces the need for real-world safety-critical testing. Potential misuse (generating adversarial pedestrian behaviors) warrants future study.
7. **Datasets:** AVA, nuScenes, and Waymo are publicly available AV datasets with established research licenses.