
Ego-Conditioned Pedestrian Body Motion Generation via Latent Diffusion with Full-Sequence Ego Conditioning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We introduce EgoPed, a system that generates realistic 3D full-body pedestrian
2 motions conditioned on vehicle ego trajectory observations from autonomous
3 driving data. Existing pedestrian motion synthesis methods rely on text descriptions
4 or action class labels, ignoring the driving context that fundamentally shapes
5 pedestrian behavior — how a person walks, stops, or reacts is strongly influenced
6 by nearby vehicle speed and proximity. We condition a Motion Latent Diffusion
7 (MLD) model on vehicle ego odometry by exposing the denoiser to the full 196-
8 timestep ego trajectory sequence (rather than a pooled summary), using a frozen ego
9 encoder that produces per-timestep conditioning tokens available to the denoiser
10 at every diffusion step. Experiments on a combined dataset built from nuScenes,
11 Waymo, and a small set of staged, highly interactive encounters recorded with our
12 sensor-equipped research vehicle (AVA) demonstrate a **40% FID improvement**
13 ($5.620 \rightarrow 3.392$) over the strongest pooled-conditioning baseline under a unified
14 evaluation pipeline, with the same denoiser capacity and the baseline evaluated at
15 its best training seed — a comparison we verify to be conservative. Ablations show
16 the design is robust: the frozen, contrastively pretrained ego encoder is critical
17 (unfreezing collapses per-condition diversity by 34%; a random frozen encoder
18 degenerates to the unconditional prior), while the specific attention mechanism is
19 not (cross-attention performs comparably to self-attention over the concatenated
20 sequence). Adding interaction-aware training-data sampling further improves
21 our best system to $\text{FID} = 2.880 \pm 0.084$, and the improvements transfer to a
22 scene-disjoint held-out test split and to evaluator-independent trajectory metrics,
23 providing a strong foundation for behavior-consistent pedestrian simulation in AV
24 settings.

25 1 Introduction

26 Generating realistic pedestrian body motions is a critical capability for autonomous vehicle (AV)
27 simulation. Simulation pipelines require diverse, plausible pedestrian behaviors — walking, turning,
28 stopping, reacting to traffic — to stress-test perception and planning systems before deployment.
29 While substantial progress has been made in text-conditioned 3D human motion generation Chen et al.
30 [2023], Tevet et al. [2023], Zhang et al. [2024], Guo et al. [2024], these systems generate motions
31 that are context-agnostic: a pedestrian crossing a street looks identical regardless of whether a vehicle
32 is approaching at 50 km/h or stationary.

33 In reality, pedestrian behavior is strongly coupled to the surrounding vehicle context. The ego vehicle’s
34 speed, trajectory, and proximity all influence whether a pedestrian will cross, wait, accelerate, or turn.
35 This *ego-conditioned* setting — generating full 3D body motion conditioned on vehicle ego trajectory

— is largely unexplored. Existing work in this space either generates 2D trajectory distributions only Mao et al. [2023], or uses ego context as scene conditioning without 3D body synthesis Feng et al. [2024].

This paper addresses ego-conditioned *full-body* pedestrian motion synthesis. We build on Motion Latent Diffusion (MLD) Chen et al. [2023], which compresses 263D HumanML3D body motion into a compact latent space and denoises within it, but replace the text conditioning with vehicle ego trajectory. The key architectural question is: *how should the ego trajectory be injected into the denoiser?*

We study two families of conditioning mechanisms:

- **Pooled conditioning (H2):** The ego trajectory is encoded by a transformer encoder and mean-pooled to a single 256D token, which is concatenated alongside the time embedding in the denoiser’s self-attention. This gives the denoiser a coarse summary of the ego trajectory but no temporal resolution.
- **Full-sequence conditioning (H4):** The ego trajectory is encoded to a full 196-token sequence, which is concatenated with the latent tokens and jointly self-attended at each denoising step. This exposes every denoising step to the full ego trajectory, enabling temporally fine-grained conditioning.

Our main finding is that full-sequence conditioning dramatically outperforms pooled conditioning. H4 achieves **FID = 3.392** versus the pooled baseline’s 5.620 under a unified evaluation pipeline — a 40% improvement with the same denoiser capacity, whose best checkpoint occurs *earlier* in training (epoch 3399 vs. 4399), so the gain does not come from additional compute. Adding interaction-aware training-data sampling improves our best system further to FID = 2.880. The full-sequence denoiser can attend to specific moments in the ego trajectory rather than a coarse average, producing motion that more closely follows the ground truth distribution.

A secondary but equally important finding concerns the ego encoder’s training regime. H4 uses a *frozen* ego encoder — pretrained to align mean-pooled ego trajectories with VAE motion latents, then fixed during diffusion training. We ablate this choice with H6, which unfreezes the encoder and allows it to co-adapt with the denoiser. H6 consistently *underperforms* H4 across all evaluated checkpoints: FID worsens by 50%, R-Precision@1 drops by 36%, and MultiModality collapses by 34%. We show that this occurs because the unfrozen encoder specializes to discriminative, per-condition representations that over-constrain generation diversity — consistent with the frozen-encoder paradigm in image generation (frozen CLIP in Stable Diffusion Rombach et al. [2022], adapters on frozen backbones in IP-Adapter Ye et al. [2023] and ControlNet Zhang et al. [2023b]).

Our contributions are:

1. The first system to generate full 3D HumanML3D body motions conditioned on vehicle ego odometry, trained on AVA, nuScenes, and Waymo data.
2. A demonstration that full-sequence temporal conditioning achieves a 40% FID improvement over single-token pooled conditioning under a unified evaluation pipeline (non-overlapping confidence intervals, 3 replications; confirmed on a scene-disjoint held-out split and on evaluator-independent trajectory metrics); ablations show the gain is robust to the attention mechanism and amplified by interaction-aware data sampling (best system FID = 2.880).
3. An ablation showing that frozen ego encoders are critical for generative diversity in ego-conditioned motion: unfreezing causes MultiModality collapse and worsens all metrics.
4. A design recommendation: frozen ego encoder + full-sequence conditioning (H4) is the correct architecture for this task.

2 Related Work

Ego-conditioned and context-conditioned motion generation. The closest works to ours generate 2D pedestrian trajectories conditioned on ego-vehicle or scene context. LED Mao et al. [2023] uses latent-space diffusion for stochastic trajectory prediction conditioned on past observations, demonstrating that latent-space diffusion substantially outperforms raw-space approaches for trajectory generation. UniTraj Feng et al. [2024] trains a unified trajectory-prediction model across multiple AV

87 datasets (nuScenes, Argoverse 2, Waymo Open Motion), demonstrating the value of multi-dataset
 88 training for agent behavior modeling — though for 2D trajectory prediction rather than 3D body
 89 synthesis. Our work differs in generating 263D full-body motion in the HumanML3D representation,
 90 enabling realistic body pose, gait, and foot contact — not just root position.

91 **Interaction-aware and retrieval-augmented motion generation.** ReMoDiffuse Zhang et al.
 92 [2023c] augments motion diffusion with retrieved nearest-neighbour motions cross-attended into the
 93 denoiser, an architectural pattern related to ours where external sequences condition the denoiser via
 94 cross-attention. InterGen Liang et al. [2024] jointly generates two-person interactions with mutual
 95 cross-attention. Our setting is a degenerate version where one agent (the ego vehicle) has a fixed,
 96 observed trajectory — enabling a simpler asymmetric design in which the observed ego trajectory
 97 conditions the denoiser as a fixed token sequence.

98 **Latent diffusion for human motion.** MLD Chen et al. [2023] is our backbone: it diffuses in the
 99 latent space of a pre-trained HumanML3D VAE, using a transformer denoiser conditioned on CLIP
 100 text embeddings. We retain the VAE, denoiser architecture, evaluation metrics, and DDIM sampling
 101 from MLD, replacing text conditioning with ego trajectory. MDM Tevet et al. [2023] operates in raw
 102 motion space and conditions on text, action labels, or trajectory waypoints. T2M-GPT Zhang et al.
 103 [2023a] uses discrete VQ-VAE tokenization — less suitable for continuous ego trajectory signals.
 104 MoMask Guo et al. [2024] achieves state-of-the-art FID on HumanML3D with masked transformer
 105 generation; we note that our metrics are computed in the same evaluation protocol, though on a
 106 different ego-conditioned dataset split.

107 **Classifier-free guidance.** CFG Ho and Salimans [2022] trains conditional and unconditional
 108 models jointly, enabling quality-diversity tradeoff at inference. We find that CFG interacts strongly
 109 with conditioning granularity: pooled conditioning (H2) shows monotonically increasing FID with
 110 CFG scale, while full-sequence conditioning (H4) is nearly CFG-insensitive. This suggests that richer
 111 conditioning naturally prevents the over-constraint that degrades FID at high CFG.

112 **Frozen encoders in generative models.** A recurring pattern in conditional image generation is that
 113 frozen encoders preserve generative diversity better than fine-tuned ones. Stable Diffusion Rombach
 114 et al. [2022] freezes CLIP during diffusion training; IP-Adapter Ye et al. [2023] and ControlNet Zhang
 115 et al. [2023b] add conditioning adapters on frozen diffusion backbones. Conversely, approaches that
 116 fine-tune the generative stack end-to-end, such as DreamBooth Ruiz et al. [2023], are known to trade
 117 diversity for condition fidelity and require explicit prior-preservation regularization to counteract it.
 118 Our H6 ablation provides direct empirical evidence for the same phenomenon in motion generation:
 119 unfreezing the ego encoder degrades generation quality and collapses per-condition diversity.

120 3 Method

121 3.1 Problem Formulation

122 Given a vehicle ego trajectory $\mathbf{e} \in \mathbb{R}^{T \times 2}$ (2D position over $T = 196$ timesteps), we aim to generate
 123 a plausible 3D pedestrian body motion $\mathbf{x} \in \mathbb{R}^{T \times 263}$ in the HumanML3D representation Guo
 124 et al. [2022]. The 263D representation encodes root velocity, root angular velocity, joint positions,
 125 velocities, and foot contact labels, fully parameterizing a SMPL-like body in motion.

126 We model the conditional distribution $p(\mathbf{x} \mid \mathbf{e})$ using a latent diffusion model that denoises within a
 127 pre-trained VAE’s latent space.

128 3.2 Motion VAE

129 We use the MLD VAE Chen et al. [2023], a Transformer encoder-decoder that compresses $T \times 263$
 130 motion to a latent $\mathbf{z} \in \mathbb{R}^{L \times 256}$ with $L = 4$ latent tokens (L is a hyperparameter; we find $L = 4$
 131 significantly outperforms $L = 8$, see Section 5). The VAE is pre-trained on the HumanML3D dataset
 132 and held fixed throughout diffusion training.




Figure 1: EgoPed architecture. The ego trajectory is encoded by a *frozen* transformer encoder into 196 per-timestep conditioning tokens. At every denoising step, the noisy latent $\mathbf{z}_t$ is concatenated with the full 196-token ego sequence (and the timestep embedding) and jointly processed by a self-attention transformer denoiser; the latent output tokens predict the noise. Diffusion runs in the latent space of a frozen HumanML3D VAE; only the denoiser is trained. Section 5.5 ablates an alternative cross-attention denoiser.

133 3.3 Ego Encoder

134 The ego trajectory $\mathbf{e} \in \mathbb{R}^{196 \times 2}$ is encoded by a lightweight transformer encoder:

$$\mathbf{c} = \text{EgoEncoder}(\mathbf{e}) \in \mathbb{R}^{196 \times 256}. \quad (1)$$

135 The encoder consists of a linear projection from 2D to 256D, followed by sinusoidal positional
 136 encodings, and $N = 4$ transformer encoder layers with 4 heads. Unlike the pooled variant (H2),
 137 which applies mean-pooling to produce a single token $\mathbf{c} \in \mathbb{R}^{1 \times 256}$, our encoder preserves the full
 138 temporal sequence.

139 **Ego encoder pretraining.** Before diffusion training, the ego encoder is pretrained with a regression-
 140 based alignment objective that distills the paired motion’s VAE latent into the ego encoding. Specif-
 141 ically, we align the mean-pooled ego encoding $\bar{\mathbf{c}} = \text{MeanPool}(\mathbf{c}) \in \mathbb{R}^{256}$ with the mean-pooled
 142 VAE posterior mean $\bar{\mathbf{z}} \in \mathbb{R}^{256}$ (computed from the paired motion) using a combined MSE and cosine
 143 objective:

$$\mathcal{L}_{\text{pretrain}} = \|\bar{\mathbf{c}} - \bar{\mathbf{z}}\|^2 + \lambda_{\cos}(1 - \cos(\bar{\mathbf{c}}, \bar{\mathbf{z}})), \quad (2)$$

144 where the cosine term encourages directional alignment in addition to magnitude. After pretraining,
 145 the encoder achieves validation $\text{MSE} \approx 1.92$, matching the pooled variant (which uses an additional
 146 2-layer projection head without improvement), confirming that the projection head adds no alignment
 147 capacity.

148 **Frozen encoder (key design choice).** During diffusion training, the ego encoder parameters
 149 are *frozen*. This preserves the encoder’s learned alignment while ensuring the generated motion
 150 distribution is not over-constrained to particular ego conditions. We ablate this choice in Section 5.

151 3.4 Diffusion Denoiser

152 We use the MLD denoiser as a self-attention transformer that operates over a single concatenated
 153 token sequence. At each denoising step, the $L = 4$ noisy latent tokens $\mathbf{z}_t$, the timestep embedding

154 $\mathbf{t}_{\text{emb}}$, and the full $T = 196$ -token ego conditioning sequence $\mathbf{c}$ are concatenated into one sequence of
 155 length $L + 1 + T = 201$ and processed by the transformer; the first L output tokens are read out as
 156 the predicted noise:

$$\hat{\epsilon} = \epsilon_{\theta}(\mathbf{z}_t, t, \mathbf{c}) = \text{Transformer}([\mathbf{z}_t, \mathbf{t}_{\text{emb}}, \mathbf{c}])_{1:L}, \quad (3)$$

157 where the transformer applies self-attention over the joint sequence and the sampler computes $\mathbf{z}_{t-1}$
 158 from $\hat{\epsilon}$. Because every ego token participates in the same self-attention as the latent tokens, each
 159 denoising step has access to the full ego trajectory at per-timestep resolution: the denoiser can attend
 160 to specific moments in the ego trajectory (e.g., a braking event or sharp turn) rather than a global
 161 average. This is the key distinction from the pooled baseline (H2), whose denoiser sees only a single
 162 mean-pooled ego token ($L + 2$ tokens in total) and therefore has no temporal resolution over the ego
 163 trajectory.

164 The added cost over the pooled baseline is modest at this scale: self-attention over 201 tokens versus
 165 $L + 2$. A natural alternative is a *cross-attention* denoiser, in which the L latent tokens act as queries
 166 and the ego sequence as keys/values, reducing the attention cost to $\mathcal{O}(L \times T)$ rather than $\mathcal{O}((L + T)^2)$.
 167 We evaluate this variant in Section 5 and find it performs *comparably*, indicating that the benefit
 168 comes from exposing the denoiser to the full ego sequence, not from the specific attention mechanism
 169 used to inject it.

170 3.5 Classifier-Free Guidance

171 We train with classifier-free guidance Ho and Salimans [2022]: for 10% of training samples, the
 172 input ego trajectory is zeroed before encoding, so the encoder’s response to the all-zero trajectory
 173 serves as the unconditional embedding (the same construction is used for the unconditional branch at
 174 inference). At inference, the guided score is:

$$\hat{\epsilon}_t = \epsilon_{\theta}(\mathbf{z}_t, \emptyset) + w \cdot (\epsilon_{\theta}(\mathbf{z}_t, \mathbf{c}) - \epsilon_{\theta}(\mathbf{z}_t, \emptyset)), \quad (4)$$

175 with guidance scale w . We find $w = 10$ optimal for H4 (Section 4).

176 3.6 Training Objective

177 The denoiser is trained with the standard DDPM noise-prediction objective:

$$\mathcal{L} = \mathbb{E}_{\mathbf{z}_0, \epsilon, t} [\|\epsilon - \epsilon_{\theta}(\mathbf{z}_t, t, \mathbf{c})\|^2]. \quad (5)$$

178 We use DDIM Song et al. [2021] for fast inference with 50 denoising steps.

179 4 Experiments

180 4.1 Dataset

181 As there are no datasets available for the problem investigated in this work, we generated a dataset
 182 for ego-conditioned motion generation from two complementary sources: existing large-scale au-
 183 tonomous driving datasets (nuScenes, Waymo), which supply scale and behavioral diversity, and a
 184 small, deliberately *staged* set of high-interaction scenes recorded with our sensor-equipped research
 185 vehicle AVA, which targets the close-range vehicle–pedestrian encounters that are rare in the public
 186 datasets. The vehicle is equipped with seven 128-layer LiDARs and eight RGB cameras for full 360°
 187 sensor coverage around the vehicle. **Figure ?? shows the vehicle and a detailed sensor setup.** To
 188 obtain scenes with meaningful interaction between the vehicle and a pedestrian, we staged **NUMBER**
 189 highly interactive scenarios. **An example scene of a pedestrian waving a vehicle through can be**
 190 **seen in Figure ??.** The recorded data is brought into nuScenes format and objects in the scene are
 191 annotated with 3D bounding boxes and orientation labels using MS3D Tsai et al. [2024, 2023].

192 With the data in nuScenes format fully annotated, we extract 3D human poses in SMPL Loper et al.
 193 [2015] format using an adapted version of OmniRe Chen et al. [2025] from all three data sources:
 194 AVA, nuScenes, and Waymo Open Dataset. We estimate poses for every pedestrian instance from all
 195 available camera images, using the projected 3D bounding box to localize the person in each view.
 196 Per-instance pose estimates from different cameras are then stitched into a single continuous sequence
 197 using the extrinsic calibration of the camera rigs. Finally, the SMPL pose sequences are converted

Table 1: Dataset composition. Each sample is a paired sequence of ego odometry (196×2) and pedestrian body motion (196×263), extracted from AVA, nuScenes, and Waymo. Train/validation splits are disjoint at the scene level; the validation split is the evaluation set used for all reported metrics.

Source	Train	Validation	Total
AVA (ours)	27	7	34
nuScenes	3,637	943	4,580
Waymo Open Dataset	5,876	1,464	7,340
Total	9,540	2,414	11,954

to the 263D HumanML3D motion representation Guo et al. [2022]. In addition to the pedestrian pose sequences, we extract the synchronized ego-vehicle trajectory from each scene’s odometry. A single automated pipeline processes every (scene, pedestrian) pair end-to-end: it loads the annotated nuScenes-format scene, extracts and stitches the per-instance SMPL body sequence, resamples both the body motion and the ego odometry to a common 20 fps grid, and expresses both streams in a *pedestrian-centric* coordinate frame (translation- and yaw-normalized to the pedestrian) so that conditioning is invariant to the absolute world pose of the scene. Each resulting sample is stored as a self-contained record containing the pedestrian’s 22 body-joint positions ($\text{ped_in_ped_frame} \in \mathbb{R}^{T \times 22 \times 3}$), the ego trajectory in the same frame ($\text{ego_in_ped_frame} \in \mathbb{R}^{T \times 3}$), and the derived 263D HumanML3D motion vector ($\text{vectors_263} \in \mathbb{R}^{T \times 263}$), where the native sequence length T varies per scene and is windowed to 196 frames at training time (see the interaction-aware cropping procedure below).

Using all three data sources we obtain **11,954** paired ego–pedestrian motion samples in total, split by scene into 9,540 training and 2,414 validation samples with no scene overlap between the two splits. Per-source counts are given in Table 1. Waymo and nuScenes dominate the sample count, while the AVA recordings contribute a small but deliberately high-interaction subset of staged vehicle–pedestrian encounters.

Each sample is a paired sequence of ego odometry $\mathbf{e} \in \mathbb{R}^{196 \times 2}$ and pedestrian body motion $\mathbf{x} \in \mathbb{R}^{196 \times 263}$, both expressed in the pedestrian-centric frame.

Interaction scoring. Rather than hard-filtering samples, we assign each scenario a scalar *interaction score* that quantifies how strongly the pedestrian and ego vehicle interact, and use it both to bias training-time sampling and to guide cropping. Let $\mathbf{r}_t, \mathbf{g}_t \in \mathbb{R}^2$ denote the ground-plane (XZ) positions of the pedestrian pelvis and the ego vehicle at frame t . The score combines three factors:

$$s = \underbrace{\sum_t \|\mathbf{r}_{t+1} - \mathbf{r}_t\|_2}_{\text{pedestrian travel}} \cdot \underbrace{\frac{1}{T} \sum_t \mathbb{I}[\|\mathbf{g}_t - \mathbf{r}_t\|_2 < 5 \text{ m}]}_{\text{proximity fraction}} \cdot \left(1 + \frac{1}{180^\circ} \sum_t |\Delta\phi_t| \right), \quad (6)$$

where $\mathbb{I}[\cdot]$ denotes the indicator function, $\phi_t = \text{atan2}(\mathbf{r}_t - \mathbf{g}_t)$ is the pedestrian’s bearing relative to the ego vehicle, and $\Delta\phi_t$ is the (wrap-corrected) inter-frame change in bearing. Intuitively, a scenario scores highly when the pedestrian moves substantially, spends many frames within 5 m of the vehicle, and sweeps a large relative bearing (e.g., crossing in front of the vehicle).

Interaction-weighted sampling. The pipeline supports drawing training samples with a `WeightedRandomSampler` whose weights are the interaction scores s (floored at 0.01 so low-interaction scenarios are still occasionally seen, then normalized to sum to the dataset size), concentrating gradient updates on the ego-relevant portion of the data distribution without discarding any samples. The VAE, the pooled baseline (H2), and our best system (EgoPed-IA) train with this weighting; the H4 reference model trains with uniform sampling, and we disentangle the two effects with a full 2×2 ablation in Section 5.2. Evaluation always uses uniform (unweighted) sampling.

Interaction-aware cropping. To bound the generation problem to a fixed horizon, we crop each scenario to $T = 196$ timesteps ($\approx 40\%$ of sequences are longer). The pipeline supports centring the crop on the *closest-approach frame* $t^* = \arg \min_t \|\mathbf{g}_t - \mathbf{r}_t\|_2$, with uniform jitter of $\pm 25\%$ of the

235 window length so the moment of closest approach is not memorized at a fixed position (the window
 236 is clamped to lie within the sequence); the alternative is a uniformly random window. The VAE, H2,
 237 and EgoPed-IA use interaction-centred crops; the H4 reference model uses random windows (ablated
 238 in Section 5.2). In both modes, ego and motion streams are cropped with a shared index so they stay
 239 temporally aligned, and scenarios shorter than 196 frames are zero-padded at the end with their true
 240 length recorded for masking.

241 4.2 Evaluation Protocol

242 We report four metrics, following the HumanML3D evaluation protocol Guo et al. [2022]:

- 243 • **FID** (Fréchet Inception Distance): computes the feature-space distance between the gen-
 244 erated and real motion distributions, using the pretrained `t2m_moveencoder` (stride-2
 245 Conv1D) and `t2m_motionencoder` (BiGRU) from Guo et al. [2022]. Lower is better.
- 246 • **R-Precision@1**: top-1 retrieval accuracy of the ego condition from batches of 32 candidates
 247 (1 positive + 31 negatives), using cosine similarity between the mean-pooled VAE latent
 248 (from generated motion) and the ego encoder output. Higher is better.
- 249 • **MultiModality (MM)**: average pairwise L2 distance between 10 generated samples for
 250 the same ego condition, measuring diversity per condition. Higher indicates more diverse
 251 generation per condition, but is only meaningful at comparable conditioning fidelity — an
 252 unconditional model maximizes MM trivially (cf. Table 3) — so MM should be read jointly
 253 with R-Precision.
- 254 • **Diversity**: average pairwise L2 distance in feature space across all generated samples,
 255 measuring population-level diversity. Values *closer to the ground-truth diversity* (5.496 on
 256 our evaluation set) are better.

257 All reported values are the mean over **3 independent replications** with 95% confidence intervals.
 258 Unless stated otherwise, metrics are computed on the full validation split (2,414 samples), which
 259 is also used for checkpoint selection, under a *unified evaluation pipeline*: identical eval-time data
 260 processing (uniform sampling, random 196-frame windows) for every system, regardless of its
 261 training-time data pipeline. Earlier internal evaluations mixed per-model eval pipelines; all headline
 262 comparisons in this paper are unified. To control for selection bias, we additionally construct a
 263 *held-out test split*: the validation scenes are partitioned into two scene-disjoint halves, one reserved
 264 purely for held-out reporting (1,190 samples from 521 scenes, never used for any model or checkpoint
 265 selection of our main system). Our best system reproduces its headline result on this held-out split
 266 (FID = 3.34 ± 0.19 vs. 3.39 ± 0.18 on the full validation split), confirming the result is not an artifact
 267 of evaluating on the selection set; the conditioning-mechanism ablation (Section 5.5) is reported on
 268 this split. Note that FID and Diversity use a fixed, pretrained evaluator shared by all systems, so they
 269 are directly comparable across architectures. R-Precision, in contrast, retrieves with each model’s
 270 *own* ego encoder embedding; cross-architecture R@1 comparisons (e.g., pooled vs. full-sequence
 271 encoders) should therefore be read as indicative rather than exact, as the retrieval space itself differs.

272 4.3 Baselines

- 273 • **H2 (Pooled, our baseline)**: EgoEncoderPooled (transformer + mean-pool $\rightarrow$ single 256D
 274 token), trans_enc denoiser (concatenated self-attention over $[\mathbf{z}, \mathbf{c}_{\text{pool}}, \mathbf{t}]$). Best at epoch
 275 4399, CFG = 5.
- 276 • **H3 (Latent-8)**: Same as H2 but with latent dimension $L = 8$ (larger VAE). Best at epoch
 277 4399, CFG = 5.

278 We additionally report four non-learned or reference baselines (Table 3) that bound the metric space.
 279 **Unconditional MLD** runs our H4 model (epoch 3399) with every ego trajectory zeroed; because
 280 CFG training dropout uses the same zeroed-trajectory construction, both guidance branches receive
 281 identical embeddings and guidance cancels exactly, yielding the model’s pure unconditional prior and
 282 isolating the contribution of ego conditioning. A *separately trained* unconditional model (identical
 283 recipe, ego dropped at every training step, 5,000 epochs, evaluated through the same ego-zeroing at
 284 guidance scale 1) reaches FID 3.24 ± 0.13 — statistically indistinguishable from the full-sequence
 285 model (3.39 ± 0.18). FID therefore does not by itself credit ego conditioning: the ego-zeroed branch

of a conditional model is merely a weak unconditional generator, and a well-trained prior matches the conditional model’s *marginal*. Repeating the control on the interaction-sampling data recipe is more striking still: an unconditional model trained that way reaches $\text{FID } 1.43 \pm 0.18$ — the best marginal of any model in this paper, better than EgoPed-IA (2.88 ± 0.08) with non-overlapping intervals — while being demonstrably ego-blind ($\text{ADE } 3.00$, behavioural separation 0.02 , R@1 at chance). Interaction sampling is therefore largely a *marginal-quality* intervention: it improves the prior ($3.24 \rightarrow 1.43$) more than it improves the conditional model ($3.39 \rightarrow 2.88$), so EgoPed-IA’s FID advantage does not survive comparison against its own prior. Evidence that the conditioning is used comes from the evaluator-independent trajectory metrics (Sec. 4.5) and the behavioral probe (Sec. 4.6), which a model without ego access cannot match. **Retrieval** returns, for each test ego trajectory, the training motion whose ego trajectory is the nearest neighbour (32-point resampled L2). **VAE interpolation** decodes a random latent interpolation of two training motions. **Trajectory+kinematic** replaces all body features with the training mean, keeping either the ground-truth root velocity (*oracle*) or nothing (*mean*), simulating a trajectory-prediction-then-fitting pipeline with no learned body model. We further add two *learned* baselines trained on the identical data recipe with the same frozen ego encoder. **Ego→motion regressor**: a deterministic sequence-to-sequence transformer regressing the 263-D features with masked MSE — the conditional-mean predictor. **Raw-space diffusion (MDM-style)**: our diffusion model operated directly in the 263-D feature space without the latent VAE (a same-capacity raw-space variant in the spirit of MDM Tevet et al. [2023], not a faithful reproduction); its best checkpoint plateaus early and never approaches latent-space quality, justifying the latent backbone. **Pretrained text-to-motion MLD**: the released HumanML3D MLD checkpoint Chen et al. [2023] generating from a text prompt instead of ego trajectory, evaluated on our held-out conditions. Because it is trained on a different corpus (mocap-quality everyday motion) its gap to our in-domain models conflates domain shift with the conditioning question; we therefore read it as evidence that off-the-shelf text-to-motion does not transfer to driving-scene pedestrians, rather than as a statement about text versus ego conditioning per se.

Information ladder: how much conditioning information is needed? To separate *that* a condition helps from *how much* information it must carry, we train the same architecture with conditioning signals of increasing information content (Table 2) and evaluate all of them with the unified pipeline. Captions are synthesized deterministically from kinematics for all 14,172 sequences, and the information content is the entropy of the resulting template distribution. The *vehicle-only* vocabulary describes only what the ego-trajectory model sees (approach/pass/wait, side, proximity, ego speed and turning, timing of closest approach). The *oracle* vocabulary additionally verbalizes the pedestrian’s own ground-truth behaviour (gait, turning, stopping); it is therefore an upper bound, not a competitor. Text models use 10% caption dropout and CLIP text features, and are evaluated at the same CFG scale as the trajectory models (a sweep over $\{1, 2.5, 5, 10\}$ is monotone for the oracle model: 6.43, 5.71, 4.14, 2.82). The vehicle-only model is evaluated on a $3\text{-epoch} \times 4\text{-CFG}$ grid (epochs 99, 3399, 4999; FID ranges from 5.12 at epoch 99 with CFG 5 to 6.67); its epoch-matched checkpoint reaches FID 5.91, ADE 2.52 and behavioural separation 0.21 — roughly two thirds of the trajectory model’s per-condition gains over the unconditional prior ($\text{ADE } 2.99 \rightarrow 2.32$, separation $0.04 \rightarrow 0.33$) at a much worse marginal, so a coarse verbal description of the vehicle carries real but incomplete ego information, and the continuous trajectory supplies the rest on both axes. Its FID-best checkpoint (epoch 99, FID 5.12) is behaviourally ego-blind (separation 0.05), reproducing within a single model the dissociation between FID and conditioning.

4.4 Main Results

Table 4 shows the definitive evaluation results. H4 achieves $\text{FID} = 3.392 \pm 0.179$ at epoch 3399 and $\text{CFG} = 10$, a **40% improvement** over the pooled baseline under the unified evaluation pipeline (H2: $\text{FID} = 5.620 \pm 0.117$ at $\text{CFG} = 5$), and EgoPed-IA improves this further to 2.880 ± 0.084 . The improvement is significant: the confidence intervals are non-overlapping (H4 CI: $[3.213, 3.571]$; H2 CI: $[6.536, 6.670]$).

Interestingly, R-Precision@1 for H4 (0.548) is *lower* than H2’s baseline (0.671). Three factors make cross-model R@1 comparisons unreliable: (1) H4 generates more diverse motions per ego condition (MM: 3.956 vs. 3.708), spreading the generated distribution in feature space and naturally reducing retrieval accuracy; (2) the R@1 metric measures alignment between generated VAE latents and ego encoder outputs — as the denoiser learns richer dynamics, this alignment metric underestimates the

Conditioning signal	Information	FID ↓	ADE ↓	Separation ↑	Diversity
None (trained unconditional, $p_{\text{drop}}=1$)	0 bits	3.24 $\pm$ 0.130	2.99	0.04	5.91
Vehicle-only text (epoch-matched, 3399)	5.17 bits	5.91 $\pm$ 0.055	2.52	0.21	5.67
Oracle text (+ ground-truth pedestrian behaviour)	6.09 bits	2.82 $\pm$ 0.027	1.90	0.93	6.20
Ego trajectory (H4)	continuous, 196×2	3.39 $\pm$ 0.179	2.32	0.33	5.79
Ego trajectory + interaction sampling (EgoPed-IA)	continuous, 196×2	2.88 $\pm$ 0.084	2.64	0.15	5.67

Table 2: Information ladder (val split, unified pipeline, 3 replications; CFG 10 for all conditional models, CFG 1 for the unconditional one). FID orders the *marginals*, not the conditioning: a well-trained unconditional model matches the trajectory model, and an oracle that is told the pedestrian’s behaviour matches EgoPed-IA. Whether a condition is actually *used* is measured in Secs. 4.5 and 4.6. ADE (m) and behavioural separation (held-out split, $K=5$) show which signals are actually *used*: the trained prior is ego-blind despite matching H4’s FID, the oracle is told the answer, and interaction sampling improves the marginal but not the per-condition fidelity (bold: best non-oracle model). The vehicle-only row is the epoch-matched checkpoint; its FID-best checkpoint (5.12) does not use the condition. The dissociation is sharpest off the ladder: an unconditional model trained with EgoPed-IA’s own interaction-sampling recipe reaches FID 1.43, better than every conditional model here, at ADE 3.00 and separation 0.02 (Table 3).

341 true conditioning fidelity; (3) R@1 retrieves with each model’s *own* ego encoder (Section 4.2), so
342 the pooled and full-sequence encoders define retrieval tasks of different intrinsic difficulty. The next
343 section settles the question with an evaluator-independent metric: on actual trajectory-following, H4
344 conditions *better* than H2, not worse.

345 4.5 Evaluator-Independent Conditioning: Trajectory ADE/FDE

346 To measure conditioning fidelity without any learned evaluator or model-specific embedding, we
347 compute root-trajectory displacement errors: for each held-out test condition we generate $K=5$
348 motions, recover the pelvis ground-plane trajectory, and compare with the ground truth over the true
349 sequence length (ADE/FDE in meters; min-of- K is standard for stochastic predictors).

350 Three controls sharpen this reading. A *separately trained* unconditional model, whose FID (3.24) is
351 indistinguishable from H4’s, is 29% worse in ADE and 26% worse in FDE: identical marginals, no
352 per-condition fidelity — the trajectory metrics see what FID cannot. An *oracle* text model that is told
353 the pedestrian’s own behaviour (Sec. 2) reaches ADE 1.90, approaching the L_2 -optimal regressor
354 while remaining generative; it bounds what a perfect description of the outcome would buy. Finally,
355 EgoPed-IA — the best model by FID — is *worse* than H4 on every per-condition metric (ADE 2.64
356 vs. 2.32; minADE₅ 1.71 vs. 1.29). Interaction-aware sampling improves the marginal but not the
357 conditioning. Its FID-selected checkpoint is trained for a third as long (epoch 1299 vs. 3399); an
358 epoch-matched checkpoint recovers part of the gap (ADE 2.58, minADE₅ 1.53) but still trails H4 on
359 every metric, so the effect is attributable to the sampling itself rather than to under-training.

360 The unconditional reference is worst on mean ADE/FDE, confirming that ego conditioning demonstra-
361 bly improves trajectory-following (2.88 $\rightarrow$ 2.32 mean ADE); its deceptively low minADE₅ reflects
362 raw sample diversity, not conditioning.

363 4.6 Behavioral Validity: Intention Distributions

364 Beyond kinematic fidelity, we probe whether generated motion couples to the ego context at the level
365 of *behavioral decisions*. We derive an automatic stopping label from root trajectories (speed below
366 0.2 m/s sustained for ≥ 1 s; heuristic thresholds, documented and spot-checked) and, for each held-
367 out condition, compare the model’s sampled stopping probability $\hat{p}(\text{stop} \mid \mathbf{e}) = \frac{1}{K} \sum_k \mathbb{I}[\text{stop}(x_k)]$
368 against the ground-truth label ($K=5$; the same samples as Section 4.5). We report the *separation*
369 (mean $\hat{p}$ on ground-truth-stopping conditions minus mean $\hat{p}$ on walking conditions), the mean
370 per-condition binary entropy of $\hat{p}$, and the Brier score.

371 EgoPed-IA couples behaviour to the ego less than half as strongly as H4 (separation 0.15 vs. 0.33)
372 and under-predicts stopping (0.14 vs. a base rate of 0.22); at a matched epoch the coupling improves
373 (0.23) while the stopping calibration worsens further (0.13), consistent with the interaction-weighted

Table 3: Reference baselines bounding the metric space (val split; same evaluator, normalization, and split lists as Table 4). All rows are mean $\pm$ 95% CI over 3 replications (the unconditional row resamples the diffusion model; the non-learned rows are reseeded end-to-end). Retrieval attains near-zero FID by copying real motions but is non-generative; trajectory+kinematic shows that following the root trajectory with a static body is catastrophic. Ego conditioning improves EgoPed’s FID by 34% over its own unconditional prior ($5.18 \rightarrow 3.39$) — and even the unconditional H4 prior outperforms conditionally generated H2 motions, underscoring the architectural advantage of full-sequence conditioning.

Method	FID $\downarrow$	Diversity	R@1 $\uparrow$
Retrieval (NN ego)*	0.062 $\pm$ 0.004	5.34	—
Trajectory+kinematic (mean)	49.60 $\pm$ 0.087	0.46	—
Trajectory+kinematic (oracle)	47.67 $\pm$ 0.085	0.72	—
Pretrained text-to-motion MLD Chen et al. [2023] [§]	14.07	3.12	— [‡]
Ego $\rightarrow$ motion regressor (learned, deterministic)	35.47	3.47	—
Raw-space diffusion (learned, MDM-style)	26.43 $\pm$ 0.098	3.55	— [‡]
VAE latent interpolation	7.91 $\pm$ 0.033	5.82	—
Unconditional EgoPed (ego zeroed)	5.18 $\pm$ 0.209	5.96	0.031 [†]
Unconditional EgoPed (trained, $p_{\text{drop}}=1$) [¶]	3.24 $\pm$ 0.130	5.91	0.031 [†]
+ interaction sampling [¶]	1.43 $\pm$ 0.183	5.76	0.031 [†]
H2 (Pooled)	5.620	5.82	0.665
H4 (EgoPed)	3.392	5.79	0.548
EgoPed-IA (ours)	2.880	5.67	0.323

*Non-generative upper bound: copies real training motions, so low FID reflects memorization rather than conditioned synthesis. [†]Exactly chance level ($1/32 = 0.031$), confirming the R-Precision protocol: without ego information, retrieval is random. [‡]R-Precision is undefined for these models (no shared embedding space with the ego encoder). [¶]Trained with $p_{\text{drop}} = 1$ (never sees the ego) and evaluated through the same ego-zeroing path as the row above. Training-time validation feeds real ego to these models and is not a valid selector, so we evaluate three late checkpoints (epochs 3999/4499/4999) offline and report the best: epoch 4999 for the standard recipe (the three span 3.24–4.90) and epoch 4499 for the interaction-sampling recipe (span 1.43–2.67; even its worst checkpoint matches EgoPed-IA). [§]Official HumanML3D checkpoint driven by a fixed text prompt; we report the *best* of three prompts (a naive scene description, a HumanML3D-idiomatic phrase, and per-sample oracle captions derived from ground-truth motion), which spanned FID 14.07–24.84. Its output is converted from HumanML3D to our normalization before scoring with the shared evaluator.

374 training distribution shifting behavioural calibration rather than with under-training. Across Tables 5
375 and 6, H4 — not the FID-best EgoPed-IA — is the best-*conditioned* model.

376 Path-crossing labels (pedestrian/ego polyline intersection) occur at only a $\sim 2\%$ base rate in our
377 windows and are omitted from the primary analysis.

378 4.7 CFG Scale Analysis

379 We sweep $\text{CFG} \in \{5, 10, 15\}$ for H4 at epoch 3399 (Table 7). A striking finding: H4 FID is nearly
380 flat across CFG scales ($3.842 \rightarrow 3.617$), while H2 shows steep monotonic FID degradation at higher
381 CFG ($6.603 \rightarrow 7.400$). This indicates that full-sequence conditioning naturally provides a rich,
382 overdetermined conditioning signal — the denoiser attends to 196 ego tokens, so higher guidance
383 amplifies an already high-fidelity signal rather than forcing mode collapse onto a coarse pooled token.

384 5 Ablation Study

385 5.1 Latent Dimension: 4 vs. 8 (H3)

386 To test whether a larger VAE latent space improves generation quality, we train a latent-8 VAE and
387 ego-conditioned diffusion model (H3). As shown in Table 4, H3 achieves $\text{FID} = 6.677$ — **19%**
388 **worse** than H2’s 5.620 at the same epoch and CFG (unified evaluation). R-Precision is essentially
389 unchanged (H3: 0.663 vs. H2: 0.665), confirming that the bottleneck for conditioning quality is the
390 ego encoder, not the latent space.

Table 4: Main evaluation results under a *unified* evaluation pipeline (identical eval-time data processing for all systems; each method at its own optimal CFG, sweep in Table 7). All values are mean $\pm$ 95% CI over 3 replications. Best values in **bold**; for Diversity, values closer to the ground-truth reference are better. EgoPed-IA additionally uses interaction-aware training-data sampling (Section 5.2); its FID replicates across two training seeds (2.880 and 2.671), whereas H4’s seed sensitivity is much larger — both are quantified in Section 5.3.

Method	Epoch	CFG	FID $\downarrow$	R@1 $\uparrow$	MM $\uparrow$	Diversity
H2 (Pooled, baseline)	4399	5	5.620 $\pm$ 0.117	0.665 $\pm$ 0.003	3.708	5.819
H3 (Latent-8 VAE)	4399	5	6.677 $\pm$ 0.135	0.663 $\pm$ 0.002	3.374 $\pm$ 0.356	5.887
H4 (Ours, full-sequence)	3399	10	3.392 $\pm$ 0.179	0.548 $\pm$ 0.002	3.956 $\pm$ 0.406	5.794
EgoPed-IA (Ours, + interaction sampling)	1299	10	2.880 $\pm$ 0.084	0.323 $\pm$ 0.004	3.371 $\pm$ 0.361	5.671
H6 (Ablation, unfrozen encoder)	3399	10	5.079 $\pm$ 0.034	0.352 $\pm$ 0.009	2.618 $\pm$ 0.309	6.064
Ground truth (reference)	—	—	—	—	—	5.496

Table 5: Root-trajectory displacement errors (meters, held-out test split, $K=5$ samples per condition, $N=1,190$). H4 beats the pooled baseline on *all four* metrics (paired bootstrap over conditions: ADE, minADE, and minFDE significant) despite its lower R@1 — direct evidence that the cross-model R@1 gap is an embedding-space artifact. The deterministic regressor minimizes expected L2 by construction yet is unrealistic (FID 35.5, Table 3); with only 5 samples H4’s minADE beats even this L2-optimal predictor. Mean ADE rewards mode concentration (collapsed H6 scores best) while min-of- K rewards coverage — the two must be read together. Bold marks the best ego-trajectory-conditioned model per column; the unconditional and oracle rows are reference bounds (min-of- K rewards the diversity of an unconditional sampler and is not a conditioning test).

Model	ADE $\downarrow$	FDE $\downarrow$	minADE ₅ $\downarrow$	minFDE ₅ $\downarrow$
H4 (full-sequence)	2.320	4.826	1.291	2.591
H2 (pooled)	2.434	4.990	1.482	2.979
H6 (unfrozen)	2.227	4.578	1.528	3.090
Regressor (deterministic, $K=1$)	1.800	3.671	—	—
EgoPed-IA (+ interaction sampling)	2.635	5.479	1.707	3.505
EgoPed-IA, epoch-matched to H4 (3399)	2.578	5.339	1.526	3.112
EgoPed-IA, seed 2 (best-val, epoch 899)	2.608	5.469	1.732	3.533
EgoPed-IA, seed 2 at epoch 1299	2.576	5.355	1.693	3.448
Unconditional (ego zeroed)	2.880	5.876	1.228	2.454
Unconditional (trained, $p_{\text{drop}}=1$)	2.990	6.093	1.202	2.380
+ interaction sampling (FID 1.43)	3.005	6.233	1.379	2.754
Vehicle-only text (epoch 3399)	2.516	5.094	1.375	2.753
Vehicle-only text, FID-best checkpoint (epoch 99)	2.709	5.527	1.312	2.568
Oracle text (told the pedestrian’s behaviour)	1.899	3.868	1.298	2.580

391 The mechanism: the fixed-capacity denoiser must now model an 8D latent distribution instead of 4D, a
392 strictly harder task, without any additional denoiser parameters. The marginal diversity improvement
393 (Diversity: 5.842 vs. 5.779, +1.1%) does not compensate. **Lesson:** increasing latent dimension is a
394 poor lever for FID improvement without scaling the denoiser capacity proportionally.

395 5.2 Interaction-Aware Sampling: a Training-Data 2 $\times$ 2

396 The interaction-aware pipeline of Section 4.1 (score-weighted sampling + closest-approach cropping)
397 is a training-*data* intervention, orthogonal to the conditioning architecture. We disentangle the two
398 with a full 2 $\times$ 2: each conditioning architecture (pooled H2, full-sequence H4) trained with and
399 without the interaction pipeline, all evaluated under the unified pipeline.

400 Two conclusions follow. First, the conditioning-granularity gap ($\approx$ 2.7–2.8 FID in both columns) is
401 roughly five times the pipeline effect, confirming that full-sequence conditioning — not data curation
402 — is the primary driver. Second, the two interventions are approximately additive: interaction-
403 aware sampling contributes a similar absolute gain to both architectures (−0.59 pooled, −0.51

Table 6: Behavioral probe on the held-out split ($N=1,190$; GT stopping base rate 0.224). The unconditional model is behaviorally ego-blind (separation ≈ 0), validating the probe. H4 couples behavior to the ego better than the pooled baseline (higher separation, best Brier) — a third evaluator-independent axis, after FID and ADE/FDE, on which full-sequence conditioning wins despite its lower R@1. H6 shows the MultiModality collapse in behavior space: the highest separation at the *lowest* entropy — it commits to a single behavior per condition — yet its Brier is no better than H2’s (overconfidence). Bold marks the best ego-trajectory-conditioned model per column. The trained unconditional model is as ego-blind as the zeroed one (separation 0.04, Brier at chance); the oracle’s 0.93 is the probe reading its own label back from the caption (“... and then stops”), which validates the probe rather than a conditioning claim.

Model	Gen. stop-rate	Separation $\uparrow$	Entropy	Brier $\downarrow$
H4 (full-sequence)	0.176	0.333	0.238	0.149
H2 (pooled)	0.226	0.285	0.297	0.178
H6 (unfrozen)	0.222	0.364	0.178	0.178
EgoPed-IA (+ interaction sampling)	0.142	0.152	0.238	0.193
EgoPed-IA, epoch-matched to H4 (3399)	0.129	0.228	0.168	0.176
EgoPed-IA, seed 2 (best-val, epoch 899)	0.090	0.138	0.157	0.189
EgoPed-IA, seed 2 at epoch 1299	0.099	0.178	0.156	0.179
Unconditional	0.194	0.021	0.533	0.199
Unconditional (trained, $p_{\text{drop}}=1$)	0.232	0.041	0.590	0.199
+ interaction sampling (FID 1.43)	0.112	0.016	0.338	0.203
Vehicle-only text (epoch 3399)	0.230	0.214	0.387	0.183
Vehicle-only text, FID-best checkpoint (epoch 99)	0.101	0.051	0.301	0.187
Oracle text (told the pedestrian’s behaviour)	0.243	0.933	0.023	0.033

Table 7: CFG sweep. H4 FID is nearly flat across guidance scales; H2 degrades steeply. H2 rows are reported under its original evaluation pipeline (the sweep predates our unified protocol); the within-model *trends*, which are the point of this table, are unaffected.

Method	CFG	FID $\downarrow$	R@1 $\uparrow$	MM $\uparrow$
H2 (Pooled)	5	6.603 ± 0.067	0.671 ± 0.003	3.503
H2 (Pooled)	10	6.963 ± 0.035	0.767 ± 0.006	3.031
H2 (Pooled)	15	7.400 ± 0.016	0.797 ± 0.005	2.724
H4 (Ours)	5	3.842 ± 0.108	0.506 ± 0.009	4.141
H4 (Ours)	10	3.392 ± 0.179	0.548 ± 0.002	3.956
H4 (Ours)	15	3.547 ± 0.122	0.536 ± 0.003	3.910

404 full-sequence), and for the full-sequence model the best checkpoint additionally arrives at 1/3 the
405 epochs. EgoPed-IA, the +pipeline full-sequence cell, is our best system.

406 5.3 Training-Seed Sensitivity

407 We retrain H4 from two additional seeds under its original recipe and select each seed’s best check-
408 point by training-time validation FID (the same protocol as the main results). Definitive best-
409 checkpoint FID across the three seeds is 3.392, 4.719, and 4.803 (mean 4.30 ± 0.79): the headline
410 seed is markedly better than the other two, so single-seed point estimates in this setting overstate
411 precision, and we report the spread explicitly. The cross-*architecture* conclusions are unaffected —
412 every H4 seed remains well below the pooled baseline (5.620), and the H2 baseline itself proved
413 training-unstable (two additional H2 seeds remained above FID 10.9 within twice the original train-
414 ing budget), meaning all our comparisons use the baseline’s *best* seed and are conservative in the
415 baseline’s favor. EgoPed-IA, in contrast, *replicates*: a second seed trained with an identical recipe
416 and selected by the same rule reaches FID 2.671 ± 0.031 (its best-validation checkpoint is epoch
417 899) against the reported seed’s 2.880 ± 0.084 , i.e. a seed mean of 2.78 ± 0.15 — a spread an order
418 of magnitude tighter than H4’s. The per-condition picture replicates as well: the second seed reaches
419 ADE 2.61 and behavioural separation 0.14 (reported seed: 2.64 and 0.15), both short of H4 (2.32
420 and 0.33), so EgoPed-IA’s weaker conditioning fidelity is a property of the interaction-sampling




Figure 2: Qualitative results from H4 (epoch 3399, CFG=10) on three ego-conditioned samples from our test set. Each row shows five timesteps of the generated 3D body motion (HumanML3D, rendered as a rigged skeleton on a small ground plane) alongside a bird’s-eye-view trajectory plot (rightmost column) showing the predicted pedestrian root trajectory (blue), ground-truth root trajectory (orange), and ego-vehicle trajectory (red). The body motion exhibits realistic gait, foot contact, and orientation changes that respond to the ego context.

Table 8: Training-data pipeline 2×2 (FID $\downarrow$, unified eval, 3 replications; best checkpoint per cell). The two interventions are approximately additive: conditioning granularity dominates (≈ 2.7 – 2.8 FID in both columns), and interaction-aware sampling contributes a similar absolute gain to both architectures (-0.59 pooled, -0.51 full-sequence).

	– interaction sampling	+ interaction sampling
H2 (pooled)	6.214 $\pm$ 0.105	5.620 $\pm$ 0.117
H4 (full-sequence)	3.392 $\pm$ 0.179	2.880 $\pm$ 0.084

421 recipe rather than of one unlucky seed. At a matched epoch (1299) the second seed gives FID
422 3.320 ± 0.139 , which shows that the two seeds reach their best marginal at different epochs and that
423 per-seed checkpoint selection, not a fixed epoch, is the meaningful protocol.

424 5.4 Encoder Freeze vs. Unfreeze (H6)

425 The most important ablation is whether to freeze the ego encoder during diffusion training. H6 uses
426 an identical architecture to H4 but with FREEZE_EGO=False, allowing the encoder to co-adapt with
427 the denoiser via backpropagation through the diffusion loss.

428 Figure 3 shows the H6 training trajectory (FID and R@1 vs. epoch). Two behaviors are notable:

- 429 • **FID/R@1 tradeoff:** As training progresses, R@1 improves (encoder specializes discrimina-
430 tively) but FID degrades (generation quality decreases). The two optima do *not* coincide:
431 training-time FID is best at epoch 2099 (4.789), while R@1 continues rising long after FID
432 has begun to deteriorate.
- 433 • **R@1 saturation far below H4:** After epoch ≈ 3299 , training-time R@1 saturates, oscillating
434 in the 0.37–0.42 band through epoch 4999 (maximum 0.417 at epoch 4599) while FID
435 remains between 5.1 and 6.3. Neither metric ever approaches H4’s operating point (FID
436 3.392, R@1 0.548) at any point in 5000 epochs.

437 Definitive evaluations (3 replications, CFG = 10) confirm the picture:

438 H6 is worse than H4 on *all metrics at all evaluated checkpoints*. Later checkpoints cannot change
439 this conclusion: even H6’s best training-time R@1 in the saturation band (0.417 at epoch 4599,
440 with training-time FID 6.339) remains far below H4’s definitive R@1 of 0.548 — and training-time
441 metrics consistently *overestimate* the definitive 3-replication values (e.g., $0.398 \rightarrow 0.348$ at epoch

Table 9: H6 definitive evaluation at three checkpoints. H4 reference included for comparison.

System	FID ↓	R@1 ↑	MM ↑	Diversity
H6 epoch=2099 (best train FID)	4.952 $\pm$ 0.088	0.263 $\pm$ 0.009	2.455	5.957
H6 epoch=3299 (best seg-1 R@1)	5.311 $\pm$ 0.060	0.348 $\pm$ 0.004	2.581	6.034
H6 epoch=3399 (matched to H4 best)	5.079 $\pm$ 0.034	0.352 $\pm$ 0.009	2.618	6.064
H4 epoch=3399 (FREEZE=True)	3.392 $\pm$0.179	0.548 $\pm$0.002	3.956	5.794




Figure 3: H6 training trajectory (FREEZE_EGO=False): training-time validation FID (left axis, red) and R@1 (right axis, blue) vs. epoch, over the full 5000-epoch run. FID is best at epoch 2099 and degrades as R@1 rises; after epoch $\approx$ 3299, R@1 saturates in the 0.37–0.42 band while FID oscillates between 5.1 and 6.3. Dashed reference lines show H4’s best operating point (FID=3.392, R@1=0.548 at epoch=3399, CFG=10). H6 never approaches H4’s performance on either metric.

3299). The MultiModality collapse is particularly striking: H6 MM = 2.618 vs. H4 MM = 3.956 (34% lower). This indicates near-deterministic generation per condition — the ego encoder has learned highly discriminative features for each trajectory, collapsing the generated motion distribution to a narrow region.

Mechanism. When the encoder is frozen (H4), it maps ego trajectories to general-purpose representations aligned with the VAE latent space. The denoiser learns to query these fixed representations and generate diverse, plausible motions for each condition. When the encoder is unfrozen (H6), backpropagation through the denoising loss pushes the encoder toward representations that *minimize the denoising loss* — essentially learning discriminative features that uniquely identify each trajectory. This over-specialization constrains generation: the denoiser outputs similar samples for each unique ego representation, collapsing MultiModality.

This pattern mirrors the frozen-encoder paradigm in image generation Rombach et al. [2022], Ye et al. [2023], Zhang et al. [2023b]: frozen conditioning encoders preserve general-purpose representations and allow the generative model to learn diverse ways to satisfy the conditioning signal.

Is the pretraining necessary? Freezing alone is not sufficient: an H4 model trained with a *randomly initialized* frozen ego encoder (identical otherwise) reaches only FID = 4.968 ± 0.171 with R-Precision@1 exactly at chance — statistically indistinguishable from the model’s unconditional prior (FID 5.18, Table 3). Together with H6, this brackets the design from both sides: the contrastive pretraining is what *injects* usable ego information, and freezing is what *preserves* generative diversity while using it.

5.5 Conditioning Mechanism: Self-Attention vs. Cross-Attention

Our denoiser injects the ego sequence by concatenating it with the latent tokens and applying self-attention over the joint sequence (Section 4). A common alternative for sequence conditioning is *cross-attention*, where the latent tokens act as queries and the ego sequence as keys and values, so the ego tokens never attend among themselves and the attention cost is $\mathcal{O}(L \times T)$ rather than $\mathcal{O}((L + T)^2)$. To test whether the specific mechanism matters, we train an otherwise-identical cross-attention denoiser (same frozen VAE and frozen ego encoder, same data, same effective batch size) and evaluate it on the held-out test split.

Table 10: Conditioning-mechanism ablation. Both denoisers consume the full 196-token ego sequence and differ only in how they attend to it. Evaluated on the held-out test split (a scene-disjoint half of the validation set), CFG = 10, mean $\pm$ 95% CI over 3 replications, best checkpoint each. The two mechanisms are statistically comparable (overlapping FID intervals); cross-attention provides no advantage.

Denoiser	Attention over ego	FID $\downarrow$	R@1 $\uparrow$	MM $\uparrow$
Self-attention (EgoPed)	concatenation	3.338 ± 0.190	0.537	4.14
Cross-attention	queries/keys-values	3.490 ± 0.059	0.501	4.61

As shown in Table 10, the self-attention and cross-attention denoisers achieve statistically comparable FID (3.338 vs. 3.490; 95% CIs [3.15, 3.53] and [3.43, 3.55] overlap), with self-attention holding a slight edge on retrieval alignment (R@1) and cross-attention generating marginally more diverse motion per condition (MM). Note that the comparison is, if anything, generous to the cross-attention variant: its best checkpoint was selected on the held-out split itself, whereas the self-attention model’s checkpoint was fixed in advance from training-time validation. We also observed that the cross-attention variant trained less stably, with a noisier FID-vs-epoch trajectory. **Takeaway:** the large improvement over pooled conditioning (Table 4) comes from exposing the denoiser to the *full* per-timestep ego sequence, not from any particular attention mechanism – either self-attention over the concatenated sequence or cross-attention works, and neither is required over the other.

5.6 Checkpoint Selection

A non-obvious finding (relevant to practitioners) is that FID does not monotonically improve with training. H2’s epoch trajectory shows FID = 6.603 at epoch 4399, FID = 8.400 at epoch 4599 (regression), and FID = 7.510 at epoch 4999 (partial recovery) — despite flat training loss throughout. H4 similarly peaks at epoch 3399 (definitive FID = 3.392) and degrades steadily thereafter (Fig. 4). **Lesson:** the final checkpoint is not the best checkpoint. Periodic FID validation during training is essential; early stopping on FID is critical.

6 Discussion

Architecture recommendation. The experimental evidence strongly favors the H4 architecture: frozen ego encoder + full-sequence conditioning denoiser (self-attention over the concatenated latent and ego sequence; cross-attention works equally well, Section 5.5). Future work seeking to improve conditioning fidelity (R@1) should explore adapter-style approaches — fine-tuning a lightweight adapter on the frozen encoder while keeping the backbone fixed — rather than end-to-end unfreezing.

R-Precision@1 and the conditioning-diversity tradeoff. H4’s R@1 (0.548) is lower than H2’s (0.671) despite dramatically better FID. This seemingly paradoxical result reflects a genuine tradeoff: richer conditioning via the full ego sequence enables the denoiser to generate *more diverse* motions per condition, spreading the generated embedding distribution and reducing retrieval accuracy. In the context of AV simulation, diversity within valid behaviors is a feature — a single ego trajectory should plausibly lead to multiple pedestrian behaviors (crossing, waiting, accelerating). R@1 measures only conditioning alignment, not generation quality; FID captures the latter more faithfully.

Limitations.




Figure 4: H4 training trajectory: training-time validation FID (left axis, red) and R@1 (right axis, blue) across the late training segment. Epoch 3399 is the best checkpoint (highest R@1, FID within noise of epoch 3199, and best definitive 3-replication FID); both metrics then degrade steadily through epoch 4999. Training loss is flat across this range — sampling quality oscillates independently of training loss.

- **Body pose quality:** Pedestrian poses are estimated (not ground-truth 3D), introducing noise in training labels. Direct 3D body annotation from datasets like JRDB-Pose or BEDLAM would improve quality.
- **Conditioning scope:** We condition only on 2D ego trajectory. Incorporating relative pedestrian-ego geometry, ego velocity profile, or map elements could improve fidelity.
- **R@1 gap:** H4 achieves R@1 = 0.548, below H2’s 0.671. The adapter-style approach described above is a natural next step.
- **Physics:** Generated motions are not physically constrained (foot sliding, penetration). Post-hoc physics cleanup could improve realism for simulation deployment.
- **Evaluator domain shift:** FID and Diversity rely on motion-feature encoders pretrained on HumanML3D, whose motion distribution differs from in-the-wild pedestrian motion; absolute FID values should be compared only within our benchmark.
- **Single benchmark, no perceptual study:** All results are on our combined AVA/nuScenes/Waymo benchmark, and we do not yet report a human perceptual study of motion realism.

Broader impact. The primary intended use of ego-conditioned pedestrian motion generation is autonomous-driving simulation: synthesizing realistic, ego-aware pedestrian behaviors enables larger-scale safety testing without exposing real pedestrians to risk, and is particularly valuable for the long tail of rare or safety-critical interactions that are unsafe or impractical to collect. We note two risks. First, generated pedestrian motions could in principle be used adversarially, e.g. to seed motion-planning policies that learn to expect only nominal behaviors and consequently fail on out-of-distribution pedestrians; we mitigate this by reporting MultiModality and diversity metrics, which expose the breadth of the generated distribution rather than only its mean. Second, training data inherit demographic biases of the source AV datasets (AVA, nuScenes, Waymo): generated motion is conditioned on the distribution of pedestrians actually observed in U.S./European urban environments, so deployment in other settings should be preceded by additional validation. Generated motion data is synthetic and contains no personally identifiable information.

7 Conclusion

We introduced EgoPed, an ego-conditioned pedestrian body motion generation system based on latent diffusion with full-sequence ego conditioning. Our key finding is that exposing the denoiser to the

full 196-token ego trajectory sequence dramatically outperforms single-token pooled conditioning: a 40% FID improvement ($5.620 \rightarrow 3.392$) under a unified evaluation pipeline, with non-overlapping confidence intervals, verified on a scene-disjoint held-out split, on evaluator-independent trajectory metrics, and conservative with respect to both training-data curation and baseline seed selection. Interaction-aware data sampling amplifies the result to $\text{FID} = 2.880$. A paired ablation reveals that frozen ego encoders are critical: unfreezing causes 34% MultiModality collapse as the encoder over-specializes to discriminative features.

These results establish a clear design recipe for ego-conditioned motion generation in AV settings: train the ego encoder contrastively on frozen VAE latents, then freeze it during diffusion training and condition the denoiser on the full ego sequence. This recipe produces state-of-the-art ego-conditioned full-body pedestrian motions across AVA, nuScenes, and Waymo, and provides a foundation for scalable, behavior-consistent pedestrian simulation.

References

- Xin Chen, Biao Jiang, Wen Liu, Zilong Huang, Bin Fu, Tao Chen, and Gang Yu. Executing your commands via motion diffusion in latent space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2212.04048.
- Ziyu Chen, Jiawei Yang, Jiahui Huang, Riccardo de Lutio, Janick Martinez Esturo, Boris Ivanovic, Or Litany, Zan Gojic, Sanja Fidler, Marco Pavone, Li Song, and Yue Wang. Omnire: Omni urban scene reconstruction. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Lan Feng, Mohammadhossein Bahari, Kaouther Messaoud Ben Amor, Éloi Zablocki, Matthieu Cord, and Alexandre Alahi. UniTraj: A unified framework for scalable vehicle trajectory prediction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2024. arXiv:2403.15098.
- Chuan Guo, Shihao Zou, Xinxin Zuo, Sen Wang, Wei Ji, Xingyu Li, and Li Cheng. Generating diverse and natural 3d human motions from text. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Chuan Guo, Yuxuan Mu, Muhammad Gohar Javed, Sen Wang, and Li Cheng. MoMask: Generative masked modeling of 3d human motions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2312.00063.
- Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- Han Liang, Wenqian Zhang, Wenxuan Li, Jingyi Yu, and Lan Xu. InterGen: Diffusion-based multi-human motion generation under complex interactions. *International Journal of Computer Vision*, 2024. arXiv:2304.05684.
- Matthew Loper, Naureen Mahmood, Javier Romero, Gerard Pons-Moll, and Michael J. Black. SMPL: A skinned multi-person linear model. *ACM Transactions on Graphics, (Proc. SIGGRAPH Asia)*, 34(6):248:1–248:16, October 2015.
- Weibo Mao, Chenxin Xu, Qi Zhu, Siheng Chen, and Yanfeng Wang. Leapfrog diffusion model for stochastic trajectory prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2303.10895.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. arXiv:2112.10752.
- Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. DreamBooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv:2208.12242.
- Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. In *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2010.02502.

- 580 Guy Tevet, Sigal Raab, Brian Gordon, Yonatan Shafir, Daniel Cohen-Or, and Amit H. Bermano.
581 Human motion diffusion model. In *International Conference on Learning Representations (ICLR)*,
582 2023. arXiv:2209.14916.
- 583 Darren Tsai, Julie Stephany Berrio, Mao Shan, Eduardo Nebot, and Stewart Worrall. Ms3d: Leverag-
584 ing multiple detectors for unsupervised domain adaptation in 3d object detection. In *2023 IEEE*
585 *26th International Conference on Intelligent Transportation Systems (ITSC)*, pages 140–147, 2023.
586 doi: 10.1109/ITSC57777.2023.10421941.
- 587 Darren Tsai, Julie Stephany Berrio, Mao Shan, Eduardo Nebot, and Stewart Worrall. Ms3d++:
588 Ensemble of experts for multi-source unsupervised domain adaptation in 3d object detection. *IEEE*
589 *Transactions on Intelligent Vehicles*, pages 1–16, 2024. doi: 10.1109/TIV.2024.3441527.
- 590 Hu Ye, Jun Zhang, Sibio Liu, Xiao Han, and Wei Yang. IP-Adapter: Text compatible image prompt
591 adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.
- 592 Jianrong Zhang, Yangsong Zhang, Xiaodong Cun, Yong Zhang, Hao Zhao, Hongtao Lu, Xi Shen,
593 and Ying Shan. T2M-GPT: Generating human motion from textual descriptions with discrete
594 representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
595 *Recognition (CVPR)*, 2023a. arXiv:2301.06052.
- 596 Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image
597 diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*
598 *(ICCV)*, 2023b. arXiv:2302.05543.
- 599 Mingyuan Zhang, Xinyi Guo, Liang Pan, Zhongang Cai, Fangzhou Hong, Huirong Li, Lei Yang, and
600 Ziwei Liu. ReMoDiffuse: Retrieval-augmented motion diffusion model. In *Proceedings of the*
601 *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023c. arXiv:2304.01116.
- 602 Mingyuan Zhang, Zhongang Cai, Liang Pan, Fangzhou Hong, Xinyi Guo, Lei Yang, and Ziwei Liu.
603 MotionDiffuse: Text-driven human motion generation with diffusion model. *IEEE Transactions*
604 *on Pattern Analysis and Machine Intelligence*, 2024. arXiv:2208.15001.

605 A Implementation Details

606 **Ego encoder.** 4-layer transformer encoder, 4 attention heads, 256D hidden dimension (1024D feed-
607 forward), sinusoidal positional encoding for $T = 196$ timesteps. Input: 2D ego position sequence,
608 projected linearly to 256D. No projection head (we verified the 2-layer MLP projection head does not
609 improve val MSE: 1.919 vs. 1.920). Pretraining: AdamW ($\text{lr} = 1 \times 10^{-4}$, weight decay 10^{-4}), 200
610 epochs, batch size 64, linear warmup followed by cosine annealing, gradient clipping at max-norm
611 1.0.

612 **Diffusion model.** MLD denoiser, trans_enc architecture (self-attention over the concatenated
613 latent, timestep, and ego tokens): 9 layers, 4 attention heads, 256D token dimension, 1024D feed-
614 forward, latent sequence length $L = 4$. The cross-attention variant of Section 5.5 uses the same
615 hyperparameters in a trans_dec configuration (latent tokens as queries, ego sequence as keys/values).
616 Diffusion: 1000 DDPM training steps (scaled-linear β schedule), DDIM inference (50 steps, $\eta = 0$).
617 CFG dropout rate: 10%. Optimizer: AdamW, $\text{lr} = 1 \times 10^{-4}$ (constant), 5000 total epochs. Batch
618 size: 128. Training hardware: NVIDIA H100 (80 GB) on an academic HPC cluster; training runs in
619 segments of up to 24h wall time.

620 **Checkpoint selection.** We select checkpoints based on validation FID (not training loss). For
621 H4, epoch 3399 is best; we confirm this with the full training trajectory and definitive 3-replication
622 evaluation. For H2, epoch 4399 is best despite continued training to epoch 5000.

623 **Evaluation.** All evaluations use TEST.REPLICATION_TIMES=3 (3 independent generation passes),
624 DIVERSITY_TIMES=300, MM_NUM_SAMPLES=100, MM_NUM_TIMES=10. Each replication regenerates
625 the full test set with fresh sampling noise; the diffusion sampler is stochastic across replications while
626 the data pipeline and evaluator are fixed (training uses seed 1234). Confidence intervals are reported
627 as $\pm$ CI (half-width of the 95% interval across replications).

628 B H6 Training Trajectory Table

Table 11: H6 (FREEZE_EGO=False) training-time validation metrics at selected checkpoint evaluation points (validation runs every 100 epochs; full curve in Figure 3). Training-time FID is best at epoch 2099 (4.789); R@1 rises until epoch ≈ 3299 , then saturates in the 0.37–0.42 band (maximum 0.417 at epoch 4599) while FID remains between 5.1 and 6.3.

Epoch	Train-time FID ↓	Train-time R@1 ↑
99	9.835	0.036
499	8.373	0.038
999	6.011	0.065
1499	5.053	0.158
1999	4.967	0.271
2099	4.789	0.273
2499	5.385	0.332
2999	5.889	0.373
3299	5.273	0.398
3399	5.141	0.380
3499	5.415	0.367
3999	6.221	0.407
4599	6.339	0.417
4999	5.561	0.408

629 C NeurIPS Paper Checklist

- 630 1. **Claims.** *Do the main claims made in the abstract and introduction accurately reflect the*
631 *paper’s contributions and scope?*
632 Answer: [Yes]. Justification: The headline claims (40% FID improvement from full-
633 sequence conditioning under a unified evaluation pipeline; further gain from interaction-
634 aware sampling; MultiModality collapse from unfreezing the ego encoder) are supported
635 by replicated evaluations with confidence intervals in Tables 4, 8, and 9, seed-sensitivity
636 reporting for both H4 and EgoPed-IA in Section 5.3, and held-out/evaluator-independent
637 verification in Sections 4.2 and 4.5.
- 638 2. **Limitations.** *Does the paper discuss the limitations of the work?*
639 Answer: [Yes]. Justification: Section 6 lists estimated (not ground-truth) pose labels,
640 restricted conditioning scope, the R@1 gap, lack of physical constraints, evaluator domain
641 shift, and the absence of a perceptual study.
- 642 3. **Theory.** *Does the paper include theoretical results?*
643 Answer: [NA]. Justification: This is an empirical systems paper; no theoretical claims are
644 made.
- 645 4. **Experimental reproducibility.** *Does the paper fully disclose the information needed to*
646 *reproduce the main experimental results?*
647 Answer: [Yes]. Justification: Architecture, training, and evaluation hyperparameters are
648 given in Section 4 and Appendix A; the dataset construction pipeline is described in Sec-
649 tion 4.1.
- 650 5. **Code and data access.** *Does the paper provide open access to data and code?*
651 Answer: [No]. Justification: Code and the derived dataset will be released upon publication;
652 the underlying nuScenes and Waymo data are publicly available under their research licenses,
653 and the AVA recordings will be released subject to anonymization requirements.
- 654 6. **Experimental settings.** *Does the paper specify all training and test details?*
655 Answer: [Yes]. Justification: Optimizer, learning rate, batch size, epochs, CFG dropout,
656 sampler settings, and checkpoint-selection protocol are in Appendix A; evaluation settings
657 (replications, metric parameters) in Section 4.2 and Appendix A.
- 658 7. **Error bars.** *Does the paper report error bars or other statistical significance information?*
659 Answer: [Yes]. Justification: All key metrics report 95% confidence intervals over 3

independent generation replications; the headline FID comparison has non-overlapping intervals.

8. **Compute resources.** *Does the paper provide sufficient information on the compute resources?*

Answer: [Yes]. Justification: Single NVIDIA H100 (80 GB) per run, ≤ 24 h wall-time segments, ~ 2 segments per diffusion model (Appendix A); evaluations take ~ 8 minutes each.

9. **Code of ethics.** *Does the research conform to the NeurIPS Code of Ethics?*

Answer: [Yes]. Justification: The work uses licensed public AV datasets and our own staged recordings with consenting participants; generated data contains no personally identifiable information.

10. **Broader impacts.** *Does the paper discuss potential positive and negative societal impacts?*

Answer: [Yes]. Justification: Section 6 discusses safety benefits for AV simulation, adversarial-use and dataset-bias risks, and mitigations.

11. **Safeguards.** *Does the paper describe safeguards for responsible release?*

Answer: [NA]. Justification: The model generates pedestrian body motion for simulation and poses no high misuse risk requiring gated release.

12. **Licenses.** *Are existing assets properly credited and licenses respected?*

Answer: [Yes]. Justification: nuScenes and Waymo are used under their research licenses and cited; the MLD backbone, HumanML3D representation and evaluators, OmniRe, MS3D, and SMPL are cited.

13. **New assets.** *Are new assets well documented?*

Answer: [Yes]. Justification: The dataset construction pipeline (annotation, pose extraction, interaction cropping) is documented in Section 4.1; documentation will accompany the release.

14. **Human subjects.** *Does the paper involve crowdsourcing or research with human subjects?*

Answer: [Yes]. Justification: Staged interaction scenarios were recorded with informed, consenting participants; recordings follow institutional data-protection procedures, and released data will be anonymized.